

The Name of the Title Is Hope

BEN TROVATO* and G.K.M. TOBIN*, Institute for Clarity in Documentation, USA

LARS THØRVÄLD, The Thørväld Group, Iceland

VALERIE BÉRANGER, Inria Paris-Rocquencourt, France

APARNA PATEL, Rajiv Gandhi University, India

HUIFEN CHAN, Tsinghua University, China

CHARLES PALMER, Palmer Research Laboratories, USA

JOHN SMITH, The Thørväld Group, Iceland

JULIUS P. KUMQUAT, The Kumquat Consortium, USA

A clear and well-documented L^AT_EX document is presented as an article formatted for publication by ACM in a conference proceedings or journal publication. Based on the “acmart” document class, this article presents and explains many of the common variations, as well as many of the formatting elements an author may use in the preparation of the documentation of their work.

CCS Concepts: • **Do Not Use This Code → Generate the Correct Terms for Your Paper**; *Generate the Correct Terms for Your Paper*; Generate the Correct Terms for Your Paper; Generate the Correct Terms for Your Paper.

Additional Key Words and Phrases: Do, Not, Use, This, Code, Put, the, Correct, Terms, for, Your, Paper

ACM Reference Format:

Ben Trovato, G.K.M. Tobin, Lars Thørväld, Valerie Béranger, Aparna Patel, Huifen Chan, Charles Palmer, John Smith, and Julius P. Kumquat. 2018. The Name of the Title Is Hope. In *Proceedings of Make sure to enter the correct conference title from your rights confirmation email (Conference acronym 'XX)*. ACM, New York, NY, USA, 29 pages. <https://doi.org/XXXXXX.XXXXXXX>

1 Introduction

ACM’s consolidated article template, introduced in 2017, provides a consistent L^AT_EX style for use across ACM publications, and incorporates accessibility and metadata-extraction functionality necessary for future Digital Library endeavors. Numerous ACM and SIG-specific L^AT_EX templates have been examined, and their unique features incorporated into this single new template.

*Both authors contributed equally to this research.

Authors’ Contact Information: Ben Trovato, trovato@corporation.com; G.K.M. Tobin, webmaster@marysville-ohio.com, Institute for Clarity in Documentation, Dublin, Ohio, USA; Lars Thørväld, The Thørväld Group, Hekla, Iceland, larst@affiliation.org; Valerie Béranger, Inria Paris-Rocquencourt, Rocquencourt, France; Aparna Patel, Rajiv Gandhi University, Doimukh, Arunachal Pradesh, India; Huifen Chan, Tsinghua University, Haidian Qu, Beijing Shi, China; Charles Palmer, Palmer Research Laboratories, San Antonio, Texas, USA, cpalmer@prl.com; John Smith, The Thørväld Group, Hekla, Iceland, jsmith@affiliation.org; Julius P. Kumquat, The Kumquat Consortium, New York, USA, jpkumquat@consortium.net.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.

Manuscript submitted to ACM

Manuscript submitted to ACM

If you are new to publishing with ACM, this document is a valuable guide to the process of preparing your work for publication. If you have published with ACM before, this document provides insight and instruction into more recent changes to the article template.

The “*acmart*” document class can be used to prepare articles for any ACM publication — conference or journal, and for any stage of publication, from review to final “camera-ready” copy, to the author’s own version, with *very* few changes to the source.

2 Template Overview

As noted in the introduction, the “*acmart*” document class can be used to prepare many different kinds of documentation — a double-anonymous initial submission of a full-length technical paper, a two-page SIGGRAPH Emerging Technologies abstract, a “camera-ready” journal article, a SIGCHI Extended Abstract, and more — all by selecting the appropriate *template style* and *template parameters*.

This document will explain the major features of the document class. For further information, the *L^AT_EX User’s Guide* is available from <https://www.acm.org/publications/proceedings-template>.

2.1 Template Styles

The primary parameter given to the “*acmart*” document class is the *template style* which corresponds to the kind of publication or SIG publishing the work. This parameter is enclosed in square brackets and is a part of the `documentclass` command:

```
\documentclass[STYLE]{acmart}
```

Journals use one of three template styles. All but three ACM journals use the `acmsmall` template style:

- `acmsmall`: The default journal template style.
- `acmlarge`: Used by JOCCH and TAP.
- `acmtog`: Used by TOG.

The majority of conference proceedings documentation will use the `acmconf` template style.

- `sigconf`: The default proceedings template style.
- `sigchi`: Used for SIGCHI conference articles.
- `sigplan`: Used for SIGPLAN conference articles.

2.2 Template Parameters

In addition to specifying the *template style* to be used in formatting your work, there are a number of *template parameters* which modify some part of the applied template style. A complete list of these parameters can be found in the *L^AT_EX User’s Guide*.

Frequently-used parameters, or combinations of parameters, include:

- `anonymous,review`: Suitable for a “double-anonymous” conference submission. Anonymizes the work and includes line numbers. Use with the `\acmSubmissionID` command to print the submission’s unique ID on each page of the work.
- `authorversion`: Produces a version of the work suitable for posting by the author.
- `screen`: Produces colored hyperlinks.

105 This document uses the following string as the first command in the source file:
106
107 \documentclass[manuscript,screen,review]{acmart}
108
109

110 3 Methodology

111 To identify and evaluate common refactorings in LangChain and LangGraph projects, we performed a series of pointed
112 queries against the GitHub REST API. Queries were directed at the search/commits endpoint using a custom python
113 script authenticated with a GitHub API token. The commit mining script is designed such that queries are made with
114 a particular combination of keywords and a maximum number of commits to return to prevent uncontrolled data
115 acquisition.

116 To prevent duplication, logic is implemented to omit commits already present in the dataset. In the event the
117 API returned a commit meeting the specified requirements, the SHA, URL, repository, and committer username were
118 appended to the CSV file. Every commit on GitHub has a unique SHA, so this was used as the primary key to differentiate
119 between distinct observations. The commit repository and author username were both derived directly from the URL.
120

121 Each query returned commits with a series of keywords present in the commit message. This ensured relevant data
122 retrieval given the commit author explicitly stated a LangChain or LangGraph refactoring as their intention. The term
123 "refactoring" was present in each query along with either "LangChain" or "LangGraph" and an associative keyword
124 such as "maintainability", "cost", "latency", or "error-handling". Our terms of interest were compiled after a rigorous
125 review of LangChain documentation and exploratory analysis of LangChain-related commits. This approach allowed
126 our team to filter to a subset of commits with a much higher likelihood of relevance.
127

128 The addition of a third keyword of interest also ensured a more diverse distribution across different contributors to
129 technical debt such that a variety of technical debt categories can be well-represented in the data. We did not require
130 any further filtering by commit date for relevance because the LangChain library was released in October of 2022. Thus,
131 the presence of the LangChain dependency satisfied any considerations for relevance in our study.
132

133 Refactorings that may have perceived benefits to technical debt but do not directly modify the implementation of
134 LangChain or LangGraph components were considered outside of the scope of this study. While such refactorings
135 may improve agentic software applications they are typically context-dependent and may not be broadly applicable
136 to LangChain implementations. By narrowing our scope to consider only direct LangChain and LangGraph code,
137 refactorings can be generalized such that developers can follow these practices in any application built using the
138 LangChain or LangGraph libraries.
139

140 During the manual examination process, phrases or claims present in the commit message were used to aid in intent
141 classification. After a review of existing patterns in our dataset, we identified seven technical debt categories: cost,
142 performance, reliability, maintainability, scalability, compatibility, and understandability. Cost refers to the monetary
143 efficiency of the system. The vast majority of projects evaluated call on external language model providers via an
144 API. The owner of this API token is then charged for each query and typically for the number of tokens processed.
145 Thus, cost as technical debt is paramount because administrators of the software can lower the monetary investment
146 associated with running the application. Typically cost can be lowered by reducing the number of input tokens present
147 in queries or by reducing the number of queries made when possible. Performance refers to improving the perceived
148 latency of an application. Technical debt associated with performance can negatively impact the user experience of an
149 application. Reliability refers to the system's ability to mitigate failures or undesired behavior in the application. In
150 LangChain and LangGraph applications this can include graceful exception-handling, proper retry logic when calling
151

external APIs, or simply producing the desired response or action. Maintainability is a technical debt category that encompasses a developer's ability to maintain or expand upon the existing application during iterative development. Scalability in this project's scope is defined as the ability of an application to scale with an increase in users or data. Applications should be designed such that they can scale without significant reductions to other debt categories such as performance and reliability. Compatibility refers to the ability of an application to integrate with different infrastructure or platforms. Agentic applications built using LangChain oftentimes require external database technologies or cloud providers. Properly designed projects should be able to seamlessly migrate between providers of such technologies. Understandability is defined as the comprehensibility of the code base itself. Issues associated with understandability can hinder the development process as members of a team may not be able to effectively read and understand existing contributions from other members of the team. Each confirmed refactoring contributes to one or more of these technical debt categories.

Upon the retrieval of relevant commits, we then required a vigorous process of manual evaluation for each commit to determine whether or not the commit contains a valid refactoring and to assign it to the proper technical debt category. This process included reading commit messages to deduce developer motivation and manually parsing through thousands of lines of code to confirm relevant changes to the code base that reduce technical debt while maintaining functionality.

Static analysis of commit diffs allowed for the omission of false positives where a commit author claimed a particular refactoring but did not yet implement such a change in the code base. Additionally, some of the commits retrieved contained the keyword "langchain" or "langgraph" but did not explicitly leverage the official langchain or langgraph libraries. In such cases, it was determined that developers were referring to langchain or langgraph patterns in their code. That is to say that while they did not use the official langchain library, they followed a similar architecture in their implementation.

Analysis of commit diffs is a meticulous process that requires familiarity LangChain abstractions and interfacing with language model APIs. Accordingly, one of our authors is a professional software engineer specializing in building agentic applications using the LangChain and LangGraph libraries. Another author is an experienced researcher specializing in refactoring and technical debt in software engineering projects.

Generally, developer intentions were derived from the commit message while the commit diff was used to validate tangible changes associated with the motivation. Pull requests were also evaluated to further explain developer intention when available. Developers frequently claim measurable latency or cost reductions in their commit messages or pull requests. These claims are difficult to validate in isolation. Instead, this study outlines recurring patterns across a diverse sample of public repositories on GitHub.

4 Results

4.1 Quantitative Analysis

The study includes 500 commits from 328 distinct open-source projects publicly available on GitHub. Of the 500 commits evaluated, our team extracted 169 confirmed refactorings representing 33.8% of the dataset. The remaining 331 (66.2%) contained keywords of interest but upon manual analysis our authors determined it either did not contain a true refactoring or contained a refactoring not associated with the LangChain and LangGraph libraries. These confirmed refactorings were then classified into 30 different distinct refactorings which we deduced from LangChain expertise and pattern recognition. Within the true positive subset, commits contain a mean of 1.73 distinct refactorings, indicating

209 that developers frequently implement multiple refactorings in tandem. Each refactoring is also assigned to one or more
 210 technical debt categories which can be seen in [Table 2](#). True refactoring commits also addressed a mean of 1.6 distinct
 211 technical debt categories. Our team compiled these refactorings such that developers leveraging the LangChain and
 212 LangGraph libraries can make these changes to enhance their code and alleviate technical debt. Given that our data is
 213 well distributed across different repositories we can validate the refactorings extracted because they occur independently
 214 across many projects. In isolation, a self-admitted refactoring is admissible. This is why our team required an empirical
 215 study to deduce repeated patterns found in the open-source community. Several refactorings that occurred once in
 216 isolation were omitted from this study as they were deemed anecdotal.
 217

218
 219
 220 *n) Message Handling Refactor.* The most common refactor among validated refactorings was the Message Handling
 221 refactoring, occurring independently 40 times in 23.5% of true positive commits. This refactor alters message processing
 222 to handle the metadata associated with LangChain message objects. Developers may bolster their implementations by
 223 using LangChain’s message objects directly without any proprietary dictionaries or by using the metadata returned by
 224 certain message objects for additional context. This refactoring in turn enhances the maintainability, reliability, and
 225 understandability of a project. A frequent error made by developers is to naively evaluate the messages returned by
 226 turns without accounting for potentially unexpected results, such as tool call data. The Message Handling refactoring
 227 typically occurred in isolation, only occurring with another distinct refactor in 9 commits representing 22.5% of Message
 228 Handling occurrences. This behavior was unique among refactorings, as others generally occurred in large overhauls of
 229 the existing source code with many distinct changes.
 230

231
 232
 233 *n) Alter State Schema to Use Add Messages Reducer.* The second most frequently repeated refactoring was the imple-
 234 mentation of LangGraph’s add messages reducer to the state schema constructor, occurring in 13% of validated commits.
 235 Though this refactoring also frequently co-occurs with other refactorings to graph state schema. The introduction of the
 236 add messages reducer co-occurs with other graph state expansions in 18 percent of it’s instances in our dataset. The add
 237 messages reducer is also frequently introduced in the same commit as the division of graphs into sub-agents, presumably
 238 because the reducer allows for concurrent message updates to state which is advantageous when multiple agents are
 239 working simultaneously on independent tasks. Such a co-occurrence is present in 27% of the commits including the add
 240 messages reducer refactoring. These results indicate a possible codependent relationship in which each refactoring
 241 benefits from the implementation of the other. Both the sub-agents and add messages reducer refactorings address the
 242 scalability of an application, so developers may commit both refactorings with the singular goal of enhancing scalability
 243 issues.
 244

245
 246 *n) LangGraph-Specific Refactorings.* Each refactoring is listed along with their frequency in the data and associated
 247 technical debts in [Table 2](#). The discovered refactoring types consist of 8 (28.5%) LangGraph-specific changes and 20
 248 (71.4%) changes that can be applied to both LangGraph and LangChain implementations. There are no alterations that
 249 apply exclusively to LangChain projects because LangGraph applications leverage primitive LangChain objects such as
 250 message types, models, or retrievers. Regardless, it is still important that we distinguish the number of true positive
 251 commits that are composed of graph or chain implementations, as these fundamentally differ in their agent orchestration.
 252 Of the 170 true positive commits, 92 (54.1%) were graph implementations using LangGraph while the remaining 78
 253 (45.9%) observations were chain or sequence implementations. Despite an equal number of queries made containing
 254 each framework and LangGraph’s relative immaturity, we observed more valid refactorings in graph implementations.
 255

261 Of the LangGraph specific refactorings, the add messages reducer refactor and division into subagents accounted for
 262 43.5% of the LangGraph specific refactorings appearing in 24 and 16 distinct observations respectively.
 263

264 *n) Performance Enhancing Refactorings.* Of the refactoring types discussed, 9 (41%) contribute to enhancing performance.
 265 Making performance the most frequently addressed technical debt in distinct refactorings. Indicating developers
 266 have different means by which to address latency in langchain applications. Additionally, performance is addressed in
 267 52 unique commits representing 31% of the true positive commits in our dataset. This makes performance the third
 268 most frequent technical debt category addressed. Of the distinct refactorings discovered, the two most commonly
 269 addressed are the Add Messages Reducer and Asynchronous Invocation. The Add Messages Reducer refactor enhances
 270 the performance of graphs by allowing for concurrent updates to state. The Asynchronous Invocation refactor has
 271 similar motivation by allowing for graphs or chains to be executed asynchronously, preventing blocking and thereby
 272 reducing latency. For instance, in a synchronous implementation when an API call is made to an external LLM provider,
 273 all other processes must wait for this action to complete before resuming execution, thus creating reducible latency.
 274

275 *n) Cost Enhancing Refactorings.* Within the true positive dataset, cost is addressed in 30 distinct commits representing
 276 17.6% of true refactors, making Cost the fifth most frequently addressed technical debt category. The Cache Recent
 277 Query Responses, Model Initialization Arguments, Context Trimming or Summarization, and Specialized Model Routing
 278 refactorings each contribute to mitigating the technical debt associated with Cost. This represents 18% of the discovered
 279 refactoring types. The Cache Recent Query Responses refactor contributes to mitigating cost by preventing unnecessary
 280 LLM calls on repeat queries thereby reducing the number of queries made to LLMs when possible. Context Trimming
 281 or Summarization can significantly reduce cost, particularly in multi-turn interactions, by reducing the number of
 282 tokens with which an LLM is called. If a language model is called with all messages from previous turns on a particular
 283 thread this causes an uncontrolled increase in the number of tokens present in each query; this is the behavior that the
 284 Context Trimming and Summarization refactor prevents. Certain model initialization arguments such as maximum
 285 iterations, maximum tokens, or maximum retries can act as guardrails preventing costly agent behavior. By setting
 286 tuning these parameters correctly in the context of the application, developers can prevent users from making calls with
 287 many tokens or prevent agents from making many LLM reasoning calls on a single user request. Thus, this refactoring
 288 type can greatly enhance the cost efficiency within an application while keeping the observable behavior relatively
 289 unaltered. The Specialized Model Routing Refactor can significantly reduce the cost of an application by routing to
 290 cheaper LLM APIs when the query is simple and only calling expensive state of the art models when a query is complex
 291 or intensive. The pricing of LLM APIs varies greatly depending upon provider, model size in parameters, and recency of
 292 the model release. As a result, routing to different models based on the perceived query difficulty is an effective way to
 293 reduce cost while maintaining the quality of agent actions.
 294

295 *n) Maintainability Enhancing Refactorings.* Maintainability is the second most frequent technical debt category
 296 addressed within the true refactoring subset. 85 commits (50%) addressed maintainability as technical debt. The
 297 refactoring types that contribute to enhancing maintainability include: ChatPromptTemplate, Message Handling,
 298 Generic Model Interface, Model Initialization Arguments, Migrate to Core Modules, Initialize tool with @tool decorator,
 299 Modularization of Components, and Graph Restructure. This technical debt category has the greatest coverage across
 300 refactoring types with 10 (45.5%) types contributing to Maintainability. The ChatPromptTemplate refactor enhances
 301 maintainability by providing one consistent and reusable prompt template throughout the entirety of a langchain or
 302 langgraph implementation. This allows for the seamless prepending of a system prompt when necessary and the prompt
 303

313 template integrates well with langchain model objects. The Message Handling refactor contributes to maintainability by
 314 improving the processing of new message objects and their associated metadata. When developers make changes that
 315 leverage LangChain's native message objects this prevents the need for manually appending and parsing a dictionary
 316 of role keys and content values. Commits leveraging the Generic Model Interface refactor type enhance maintainability
 317 by allowing developers to seamlessly switch between models by changing a single value within the model object. If a
 318 developer uses a provider-specific model object then opting to a different model or provider will require alterations to
 319 the model arguments and the objects with which the model interacts. These changes can quickly become expansive
 320 throughout the code base. By using a universally compatible model object we can prevent unnecessary changes when
 321 switching between models and improve the maintainability of an application. The Initialize Tools with @tool Decorator
 322 refactor contributes to maintainability by defining tools and their functionality simultaneously within the code base.
 323 The tool decorator does so by packaging the tool function as a LangChain tool object and using the function's docstring
 324 as the tool description to be provided to an agent to improve tool selection. The alternative would be to define the
 325 tool object and tool functionality separately which may make it difficult for developers to maintain the application as
 326 they may need to make changes in two distinct locations for a single tool. The modularization of components refactor
 327 contributes solely to maintainability as technical debt. In this refactor developers split distinct langchain or langgraph
 328 components into separate files which allows them to easily locate components within the file tree of an application.

329 n) *Reliability Enhancing Refactorings*. Reliability was the most frequently addressed technical debt category in the
 330 true refactorings subset. 90 of the 170 observed refactorings contributed to reliability as technical debt. Of the 22 distinct
 331 refactoring types, 8 contribute to improving reliability. These include: ChatPromptTemplate, Message Handling, Model
 332 Initialization Arguments, Compile with Checkpointer, Structured Tool Wrapper, @tool decorator, model fallbacks,
 333 and specialized model routing. The Model Initialization Arguments refactor type contributes to reliability because
 334 arguments such as maximum iterations and maximum retries improve the resilience of agent execution. The Compile
 335 with Checkpointer refactor contributes to reliability by providing a snapshot of agent state and execution to fall back to
 336 upon failure. Without the implementation of a checkpointer, the runtime state and execution history are all maintained
 337 solely within memory, so if the application crashes, a user will lose all of the data from their current thread. By saving
 338 this data to disk incrementally developers can provide a checkpoint to fall back to allowing the application to fail
 339 gracefully. LangChain's structured tool object allows developers to specify input/output schema when defining a
 340 tool. This can enhance the reliability of tools with complex inputs or outputs. The @tool decorator allows for similar
 341 assurance of inputs and outputs in less complex use cases. The tool decorator uses the tool function's docstring as a tool
 342 description providing agents with more context surrounding the tool's intended use-case. As a result, models are less
 343 likely to misuse or incorrectly select tools for a given user query. Additionally, the tool decorator uses the function's
 344 parameter and return type hints in the function signature to prevent type errors associated with agent tool use. The
 345 model fallbacks refactor is directed solely at reliability. Particularly when agentic applications are calling external APIs
 346 for model use, this refactor serves as an important mechanism making the application resilient in the event a certain
 347 model API or provider fails to respond to requests. The Specialized Model Routing refactor can optimize the tradeoff
 348 between cost and reliability by routing to state of the art models when queries are determined to be complex or crucial.
 349 This can ensure accurate model results while mitigating cost when possible. As a result, the Specialized Model Routing
 350 refactor impacts both cost and reliability.

351 n) *Scalability Enhancing Refactorings*. Scalability was addressed in 52 (30.6%) true refactorings making it the third
 352 most frequently addressed technical debt category. As depicted in Table 2 the distinct refactoring types affecting

scalability include: ChatPromptTemplate, Alter State Schema to use add messages reducer, Reorganize Agent into Hierarchical Sub-Agents, Model Initialization Arguments, Compile with Checkpointer, and Minimize State Schema. The use of an add messages reducer in state enhances scalability by coordinating seamless, concurrent updates to state. This improves scalability by allowing multiple nodes running in parallel to write new data or metadata to the shared state without blocking one another. The Reorganize Agent into Hierarchical Sub-Agents refactor contributes to scalability by dividing a singular "God" agent into a supervisor(s) and single-responsibility sub-agents. This can help a system to scale with an increase in data or queries on a single thread because data context is more isolated depending on the specific task requested. This optimization is limited in the event that every request falls to the same sub-agent and the supervisor or delegator node may become a potential bottleneck depending on the configuration. The Minimize State Schema refactor improves the scalability of a langgraph application by reducing the amount of data maintained throughout graph execution. When state schema contains unnecessary fields the data that is recorded with each turn can quickly grow in multi-turn threads and make it difficult to scale an instance over many turns for a user. Implementing a checkpointer is essential to running a langgraph application at scale. LangGraph check pointers are designed to save state into a thread as defined by a unique thread id, this in turn allows for different users or conversations to occur simultaneously without interfering with one another. Additionally, a checkpointer enables scaling across multiple workers such that any worker can pick up a request for a given thread by simply loading the most recent checkpoint and continuing to process a user request from that state. A checkpointer is also essential for human-in-the-loop workflows that pause and wait for further user input before proceeding with graph execution.

n) *Understandability*. Understandability is addressed by 21 commits (12.4%) in the true refactoring dataset. Among these, *Graph Restructure* is the most common refactoring type contributing to enhanced understandability, appearing in 8 of the 21 commits related to understandability. Restructuring a state graph to eliminate unnecessary conditional routing and minimize complexity improves understandability by making simpler and more comprehensible execution flows. In the event that a graph is large, complex, or contains many conditional edges it may take developers a great deal of effort to reconstruct the graph externally for maintainability or documentation purposes. Redesigning graphs to be more minimalistic can reduce the cognitive load associated with understanding the source code of a langgraph application. The other refactoring types contributing to understandability include: Consistent Node Naming Conventions, Consistent Tool Naming Conventions, and Modularization of Components. Among these Consistent Tool Naming Conventions is the least frequent refactoring type occurring only once in the true refactoring dataset. It is worth noting that most commits have consistent tool naming conventions prior to the changes. Proper node naming conventions can also enhance the understandability of a langgraph implementation by being self-descriptive and action-oriented such that developers reading the source code can immediately understand a node's purpose upon viewing it. The Modularization of Components refactor contributes meaningfully to understandability because it allows developers to easily locate and distinguish different components based on the file in which they are contained. A meaningfully organized file structure can make the purpose and classification of a component immediately apparent.

n) *Technical Debt Co-Occurrences*. Within true positive commits, however, reliability and maintainability were the most frequent occurring 68 and 64 times respectively. However, it is important to note that these two technical debts co-occur with one another 25 times in the true positive dataset. This means that 36.7% of reliability and 39% of maintainability occurrences address both of these technical debts simultaneously. This is because improved reliability and maintainability are often a by-product of refactorings whose explicit motivation is to address one of the two. It is also important to recognize that many of the identified refactorings address technical debt across multiple categories.

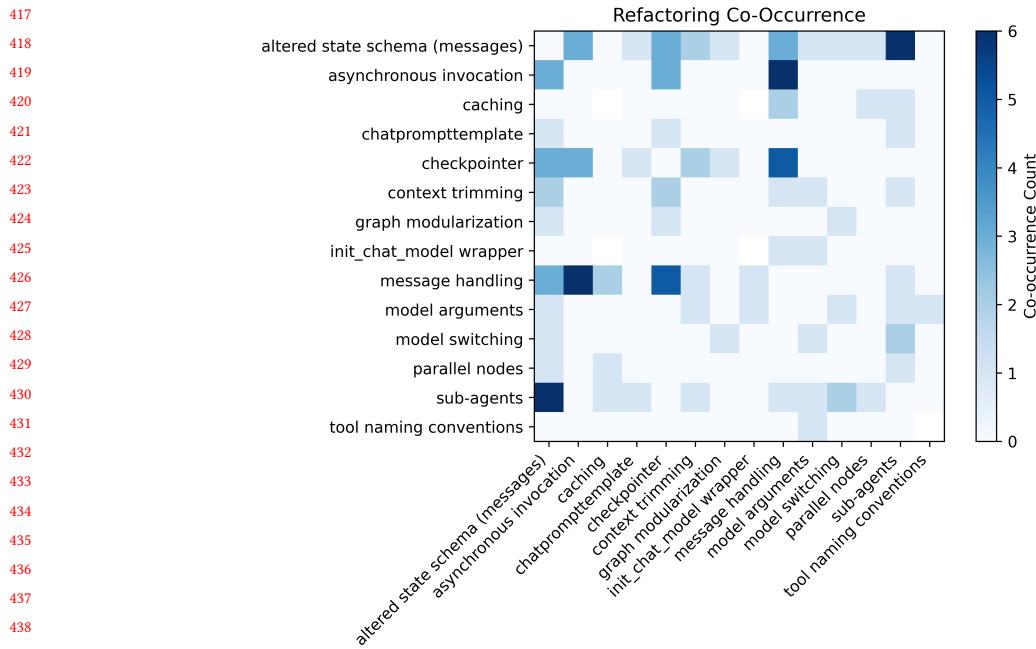


Fig. 1. Refactoring Co-Occurrence Heatmap

For instance the use of a prompt template improves both reliability and maintainability by improving language model responses and making prompt reuse seamless throughout application source code seamless for developers. Model initialization arguments are another versatile refactoring that can effectively impact the reliability, scalability, cost, performance, and compatibility of an application by setting explicit guardrails for language models and allowing for provider-specific keyword arguments.

n) Refactoring Type Co-Occurrences. Certain refactoring types frequently occur with others in the same commit. This happens, in part, because certain refactors mutually benefit from one another. The addition of the add messages reducer refactor and the hierarchical sub-agents refactor co-occur 6 times in the true positive dataset. This relationship exists because the add messages reducer function allows for concurrent state updates. When an agent is subdivided into multiple sub-agents the resulting graph will now benefit from concurrent updates in the event multiple agents are executing in parallel. Previously, the graph may have followed a simple sequential architecture without the need for complex state management; however, multiple agent nodes necessitate concurrent state management. A similar relationship exists between the Asynchronous Invocation and Message Handling refactoring types which occur in tandem 6 times as well. These refactoring types seem to have two distinct unrelated motivations and enhancements and yet developers are compelled to implement them within the same commit at a relatively high rate. The Message Handling and Checkpointer refactoring types have similar indicators of a relationship, occurring together 5 times in the dataset. This relationship is intuitive because as developers are implementing a checkpointer to improve memory persistence this refactor provides visibility to the quality of message management. That is to say that the contents of

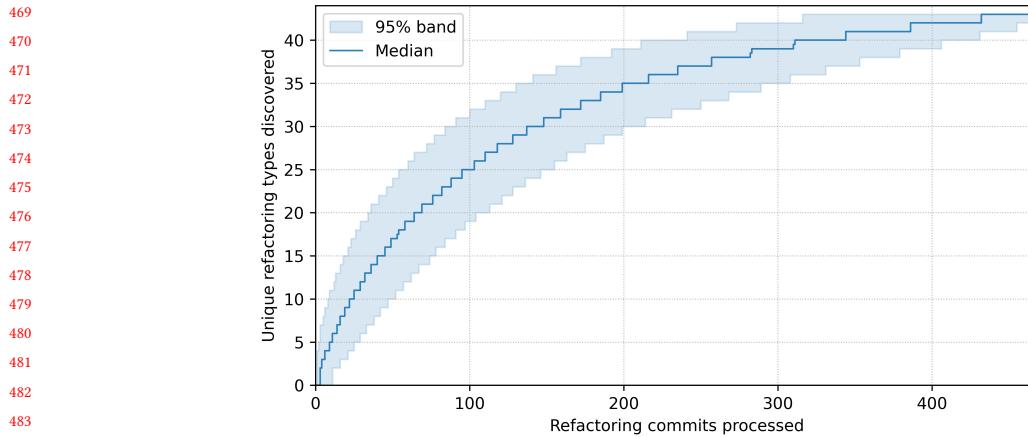


Fig. 2. Coverage of the Refactoring-Type Catalog as Data Scale Increases

memory becomes a concern when improving the memory implementation. Such changes have generally been observed in complete overhauls of memory management systems within the langchain portion of applications.

n) LangChain-Specific Refactorings. Table 3 depicts each extracted refactoring type along with the component affected and a short description of how the change is implemented. Each refactoring type affects one of eight distinct langchain components. The chain and graph component were grouped together because they each serve as the primary route for smaller interconnected components and vary only based on the scale or objective of the application. The messages component refers to message history across turns including system, human, and ai generated message objects. Messages refers specifically to conversational history or the management of message objects; this differs from other components such as state or memory. State is another distinct component that refers to the meta data passed between nodes in a langgraph state graph. The contents of state is entirely dependent upon how a developer chooses to implement it, but it is generally best practice to construct state as a typed dictionary. Model refers to the language model object which is invoked by other langchain components (e.g, nodes) typically with context such as previous messages and metadata. Rather than interacting with LLM provider APIs directly, developers generally use langchain’s prebuilt model objects to leverage the abstraction they provide. The imports component refers to the exact langchain modules developers choose to import. There is some degree of overlap in the functionality offered by different langchain modules and as such developers can achieve the same functionality with different modules. Tools the most frequently impacted component across distinct refactoring types appearing in 5 of the 30 refactoring types (16.7%). Tools are langchain objects that provide models with functionality or allow a model to interact with other existing systems. Tools vary greatly depending upon the implementation but the same underlying refactorings occur throughout repositories. Individual nodes are a component impacted in only two refactorings and this is the only component that is framework specific, existing only in langgraph implementations.

n) Diversity of Refactoring Types Over Data Acquisition. The process of data acquisition and manual labeling of commits from various projects is intensive. One might argue that with an increase in our sample size, more unique refactoring types could be discovered. Figure 2 displays the justification of our sample size. This figure displays the

521 median and 95% confidence interval of refactoring type discovery over the number of commits processed when shuffling
 522 and resampling the dataset 2000 times. This indicates that regardless of order the number of unique refactoring types
 523 discovered per new commit processed greatly decreases as the number of commits processed increases. Visually, this is
 524 represented by a significant decrease in slope as x increases. Thus, the evidence suggests that our team has extracted
 525 the majority of unique refactoring types present in open source langchain repositories.
 526

Table 1. LangChain-Specific Technical Debts

Technical Debt	Goal	Situation	Consequence	Intent Groups
Cost	Minimize the monetary cost associated with LLM API calls	LLM calls are made too frequently or with many input tokens	User interaction is expensive due to over-reliance on API calls	Reduce input tokens; Avoid iterating uncontrollably; Prevent unnecessary API calls
Performance	Minimize the perceived latency	Graph or chain takes excessive time to respond/act when prompted	User experience is negatively impacted	Maximize cache reuse; Execute components asynchronously; Prevent unnecessary API calls; Enable parallelism
Maintainability	Reduce future change effort and change-ripple while preserving behavior	Evolution tasks (e.g., updating bases/packages, changing defaults, parameterizing build/run options) currently require repeated or fragile edits across artifacts	cascading edits, higher regression risk, longer lead time for change	Divide components into isolated files; Leverage wrappers to avoid provider-specific configurations; Use LangChain message-handling abstractions;
Reliability	Mitigate errors and undesired behavior	Errors, confusion and unexpected behavior	Errors and unexpected or undesirable behavior	Enhance type safety; Expand agent context; Increase agent determinism
Scalability	Seamlessly scale with an increase in users or data	The system cannot sustain an increase in users or data	Increase in users or data causes errors, latency, or expensive operations	Limit context/state in multi-turn threads; Divide graph into sub-agents or sub-graphs
Compatibility	Integrate seamlessly with new providers and dependencies	Updating dependencies, platform versions, or runtime environments causes failures due to hardcoded assumptions or environment-specific configurations.	Changes in dependencies or platforms causes errors or undesired behavior	
Understandability	Readable and comprehensible code for iterative development	Ambiguous naming conventions, monolithic code, or convoluted structure make it difficult for developers to interpret the code base.	Reading and understanding the project source code requires additional time and effort.	Consistent component naming conventions; Divide components into isolated files; Include type hints; Include descriptive doc-strings

Table 2. Mapping LangChain Refactoring Types to Technical Debt Categories.

Refactoring Type	Commit Count	Cost	Performance	Reliability	Scalability	Maintainability	Understandability	Compatibility
Asynchronous Invocation	10		×					
Adopt ChatPromptTemplate	6			×	×	×	×	
Alter State to Use Messages Reducer	24		×		×			
Decompose Agent into Sub-Agents	16				×			
Message Handling	40			×		×	×	
Initialize Model via Universal Interface	4					×		
Cache Recent Query Responses	8	×	×					
Add Model Arguments	10			×	×	×		×
Compile Graph with Checkpointer	24			×	×			
Migrate to Core APIs	4					×		
Minimize State Schema	3		×		×			
Migrate to StructuredTool	4			×				
Initialize Tools with Decorator	4			×			×	
Standardize Node Naming	2						×	
Standardize Tool Naming	2						×	
Context Trimming or Summarization	7	×	×					
Execute tool calls in parallel	2		×					
Execute nodes in parallel	5		×					
Model Fallbacks	2			×				
Add Model Routing Logic	5	×	×	×				
Modularization of Components	6					×	×	
Restructure Graph Topology	8		×				×	

4.2 Qualitative Analysis

The qualitative analysis of this study aims to analyze some valuable examples of commits that express the refactorings discussed in the previous section. While it is valuable to note that the refactorings presented are clear repeated patterns across open-source repositories, it is still vital that we evaluate what changed in the source code before and after a valid refactoring. This section provides brief static analysis for a variety of commits from the different categories mentioned previously. Here we provide more depth as to what changed before and after given refactorings as well as why developers claim it alleviates technical debt or enhances the code base.

Example 1 (Maintainability + Scalability + Performance - Use an annotated list with add messages reducer in graph state): Developers have a great deal of autonomy in designing the state schema for their state graph in a LangGraph application. The refactoring displayed in Listing 1 aims to follow the best practice for maintaining conversational memory in the graph state per official LangGraph documentation. By creating an annotated list of BaseMessage objects with the add messages reducer, developers take advantage of the abstraction provided by LangGraph objects while simultaneously allowing for concurrent updates to the list of BaseMessage objects. By altering the state schema to follow

Table 3. LangChain-Specific Refactorings

Refactoring Type	Component	Description
Asynchronous Invocation	Graph/Chain	Graphs or Chains/Sequences are invoked asynchronously using the <code>ainvoke()</code> , <code>astream()</code> , or <code>abatch()</code> method to prevent blocking
ChatPromptTemplate	Messages	Change prompt engineering to leverage <code>ChatPromptTemplate</code> rather than manually appending dictionary
Add Messages Reducer	State	Maintain conversational history with an annotated list of <code>BaseMessage</code> objects and corresponding <code>add_messages</code> reducer
Sub-Agents	Graph/Chain	Reconstruct graph into specialized sub-agents with supervisor/delegate node(s)
Message Handling	Messages	Alter message handling to leverage <code>BaseMessage</code> objects and their associated metadata
Initialize Model with Generic Interface	Model	Initialize model with generic <code>init_chat_model</code> rather than proprietary model wrapper
Cache Recent Query Responses	Messages	Cache recent queries to reduce latency & cost on repeat queries
Robust Model Initialization Arguments	Model	Initialize model with max iterations, max tokens, timeout, temperature, top-p, and model fallbacks as necessary
Compile Graph with Checkpointer	Graph/Chain	Initialize a memory saver & unique thread id then compile your state-graph with the checkpointer
Migrate off of community modules	Imports	Migrate off of community modules when identical functionality is available in official modules
Minimize State Schema	State	Remove unused or unnecessary fields from state to reduce the amount of data transmitted between nodes
Structured Tool	Tools	Alter tool definitions to use <code>StructuredTool</code> for more robust input/output schema definitions
@tool decorator	Tools	Alter tool definitions to use the <code>@tool</code> decorator along with type hints and a descriptive docstring
Consistent Node Naming Conventions	Nodes	Alter node definitions to follow a clear & consistent naming convention in which the node name corresponds to the function name
Consistent Tool Naming Conventions	Tool	Alter tool definitions to follow clear & consistent naming conventions
Context Trimming & Summarization	Messages	When passing previous conversation turns to model trim down to the last n messages or summarize
Parallel Tool Calls	Tools	When a single query requires multiple tool calls, execute them in parallel
Parallel Nodes	Nodes	Reconstruct graph to allow for parallel execution of independent nodes
Specialized Model Routing	Model	Tailor model selection based on query analysis to ensure reliable results or minimize cost
Alter node returns to only updated fields	Node	Alter node functions to return only updated fields of state rather than the entire state
Use MCP for Tools	Tools	Use MCP as a tool service rather than LangChain tool objects
Migrate to Runnable Sequence	Chain	Migrate from chain functions to Runnable Sequence
Streaming	Graph/Chain	Stream graph/chain rather than invoking so tokens arrive as they are produced

this pattern developers aren't burdened by manually appending conversational history to a dictionary throughout the execution of the graph. Additionally, synchronous state updates can introduce some degree of reducible latency if

```

677      Listing 1. Commit ad59db5 from data-cleaning-agent: Use Annotated types for concurrent message updates
678      @@ -37,8 +38,8 @@ class DataCleaningState(TypedDict):
679          pending_tasks: List[str]
680          progress_percentage: float
681
682          # Agent communication
683          - messages: List[BaseMessage]
684          + # Agent communication - use add_messages for concurrent updates
685          + messages: Annotated[List[BaseMessage], add_messages]
686          agent_results: Dict[str, Dict]
687          communication_log: List[Dict]
688
689      multiple tools, nodes, or agents are triggered simultaneously as state updates are required throughout the execution
690      flow of the graph. Concurrent updates to the message state allow for a more seamless execution flow throughout the
691      graph that minimizes blocking and allows for components to make concurrent progress without compromising state
692      integrity.
693
694
695      Listing 2. Commit 84365f4 from hikizan-emacs: Integrated 'langgraph.checkpoint.memory.MemorySaver' for persistent agent state
696      across runs.
697      @@ -321,9 +418,18 @@ def __init__(self):
698          google_api_key=GOOGLE_API_KEY,
699          )
700
701          self.memory = MemorySaver()
702
703          self.tools = [execute_elisp_code, read_file, write_to_file, list_files, grep, find_files]
704          self.tools_by_name = {tool.name: tool for tool in self.tools}
705          - self.graph = self._build_graph()
706          +
707          # Build graph with memory checkpoint
708          + self.graph = self._build_graph().compile(checkpointer=self.memory)
709
710
711      Example n (Reliability + Scalability - Compile with checkpoint): The Listing 2 above exemplifies the compile with
712      checkpoint refactoring. This refactoring can be performed with minimal changes to the code base. Developers can
713      simply instantiate a LangGraph memory saver object and pass this memory saver as the checkpointer when compiling
714      the graph into a runnable object. Previously, the graph was compiled with no functional checkpointer, so there exists
715      no other version of state and runtime configuration to revert back to in the event of an error. For this reason, the
716      checkpointer refactoring is classified as a contribution to reliability. Additionally, a checkpointer benefits the scalability
717      of a langgraph application in multiple facets. An available persistent checkpointer allows for the data generated through
718      graph execution to be offloaded to disk when necessary rather than relying solely on state, which exists in main memory,
719      to capture execution and conversational history. This enhancement allows for a LangGraph application to reasonably
720      scale with an increase in data, such as in a long multi-turn interaction.
721
722      While certain aspects of execution history such as prior messages can be maintained in the graph state, this state is
723      maintained entirely in main memory. Thus, if a particular instance of the LangGraph application were to terminate, all
724      of its history would be lost. This is generally not practical for the deployment of software at scale. This is why the
725      LangGraph SDK offers a checkpointer which can be easily instantiated and compiled with a graph. The introduction of
726      a checkpointer and thread id allows for memory persistence across threads. This means that for a given thread there
727
728

```

Manuscript submitted to ACM

729 are a series of checkpoints that save the graph state and runtime configuration of the agentic system at a given point in
 730 time. This in turn enhances the scalability of a LangGraph system as it allows for multi-tenant applications where each
 731 user is the owner of one or more threads. Additionally, by compiling each instance of the state graph with a unique
 732 checkpointer and thread id this allows for persistent memory in the event the application crashes or loses connection.
 733

734
 735
 736 Listing 3. Commit c72917f from cm3070-lawtime: Switch to asynchronous graph invocation
 737 diff --git a/server/services/task_service.py b/server/services/task_service.py
 738 index dba2fb6..d1bac37 100644
 739 --- a/server/services/task_service.py
 740 +++ b/server/services/task_service.py
 741 @@ -139,7 +139,7 @@ async def propose_tasks(
 742 try:
 743 # Invoke the LangGraph agent - it will handle all validation and initialization
 744 logger.info(f"Invoking LangGraph agent with source_type: {source_type}")
 745 - final_state = graph.invoke(initial_state)
 746 + final_state = await graph.ainvoke(initial_state)
 747
 748 # Extract proposed tasks from final state
 749 proposed_tasks = final_state.get("proposed_tasks", [])
 750

751 Example n (Performance - Use asynchronous version of invoke): Both LangChain and LangGraph offer asynchronous
 752 versions of the stream, batch, and invoke methods. The invoke method is by far the most common method for calling a
 753 compiled state graph or chain. Calling a graph or chain asynchronously can benefit the execution of an agentic system
 754 and prevent blocking. If chains or graphs are invoked synchronously this means that when a latent task such as making
 755 an external LLM call is performed, the rest of the program is blocked and cannot make progress on other routines.
 756 Contrarily, when a graph or chain is invoked asynchronously, this allows the program to execute concurrently and
 757 in turn reduce latency by making progress on tasks while waiting on latent tasks such as LLM calls. Asynchronous
 758 methods in the LangChain ecosystem still require that node and tool functions are defined asynchronously such that
 759 they can be called as a co-routine.
 760

761
 762
 763 Listing 4. Commit f4877dc from boss-bot: Enhanced tool definitions with the '@tool' decorator, improving type safety and documentation

```
764 + """Collection of media inspection tools using modern @tool decorator"""
 765
 766 @staticmethod
 767 + @tool
 768 async def extract_video_metadata(video_path: str) -> Dict[str, Any]:
 769 -     """Extract comprehensive video metadata"""
 770 +     """Extract comprehensive video metadata using ffprobe or similar tools
 771 +
 772 +     Args:
 773 +         video_path: Path to the video file to analyze
 774 +
 775 +     Returns:
 776 +         Dictionary containing video metadata including duration, resolution, codecs, etc.
 777 +     """
 778
 779
 780
```

Example n (Reliability + Maintainability + Understandability - Use the @tool decorator with type hints): The LangChain SDK supports a tool decorator that allows for developers to seamlessly define tools while simultaneously providing informative context for a model when calling upon it. When a model calls a tool that is defined simply as a function in the codebase, oftentimes these tools require a variety of parameters to be passed of particular types. If a model passes the incorrect data type as a particular argument to a tool function this will lead to a type error and the tool call will not be able to execute. The tool decorator not only defines the function as a tool object, but also uses the function signature to provide the model with type hints to improve the reliability of tool calls. Type hints included in the function signature for both parameter and return types will mitigate type errors associated with tool calls. Additionally, the docstring is used as a tool description which is also available to the LLM. Informative docstrings can improve agent tool selection such that the most desirable tools are selected given the context and query.

Example n: Scalability + Maintainability + Reliability: Subdivide agent into specialized sub-agents. In large LangGraph applications, it can be preferable to subdivide a singular 'God agent' into single-responsibility sub-agents with a supervisor or routing node that acts as a delegator. Sub-agents can have a system prompt tailored to their granular use case as well as a tool subset that negates tools outside of their scope. Delegating to specialized sub-agents or sub-graphs can also allow for the parallel execution of independent tasks. In the event that the router or supervisor node elects multiple tasks for a given query, these co-routines can run concurrently and return to the supervisor without blocking one another.

Listing 5 portrays some key differences before and after the subdivision into modular sub-agents. As can be seen in the diff, the original version of the state graph routed to simple tools based on the presence of hard-coded keywords. This is a fragile practice with very little flexibility to respond to user interaction. There are many edge cases in which a user may adequately articulate the task which they wish to accomplish and this original logic may still fail to produce the desired results. Given the ability of modern language models to capture query intention, it is optimal to allow an agent to reason about the user query, develop a course of action, and delegate the necessary tasks. The additions shown in listing 5 show that rather than routing manually on the basis of hard-coded keywords, we can instantiate a deep agent with access to a list of sub-agents for the purpose of delegation. Sub-agents are equipped with their own unique subset of tools and task-specific system prompt. This reduction in agent scope can enhance the reliability of agent actions and reduce the likelihood of context-overload for a particular model instance. The implementation of a hierarchical sub-agent architecture such as that depicted in Listing 5 allows for seamless expansion of system functionality with minimal changes to the main workflow. Sub-agents are isolated and modular, thus, to add new agent functionality a new sub-agent can be easily created in isolation and simply added to the list of sub-agents available to the supervisor node. This enhanced ability to improve or expand workflow functionality improves the maintainability of the LangGraph implementation.

Example n: Reliability - Reduce temperature to increase model determinism. Listing 6 portrays an example of the temperature lowering refactoring. In this commit, the developer lowered the temperature value from 0.7 to 0.4 during model initialization. This is a common step taken to increase the determinism of agents. LLMs are not typically set to a default temperature of zero because this would indicate that the next generated token is always the most probable. Additionally, the developer chose to include a top-p value of 0.7 and decreased the max-tokens value by 50 percent. The top-p argument serves a similar purpose to temperature, in that it increases the determinism of model generation. Top-p (top probability) reduces the available sample size of a language model's probability distribution to only the tokens with the largest corresponding probabilities whose cumulative sum is equal to the top probability value. This

```

833 Listing 5. Commit cb3467c from hacs-ai: Complete replacement of basic graph with deepagents framework; 5 specialized admin
834 sub-agents for different domains
835 diff --git a/examples/hacs_developer_agent/deep_agent.py b/examples/hacs_developer_agent/deep_agent.py
836 new file mode 100644
837 index 0000000..08a2217
838 --- /dev/null
839 +++ b/examples/hacs_developer_agent/deep_agent.py
840 @@ -0,0 +1,345 @@
841 +def create_hacs_deep_admin_agent(
842 +    instructions: str,
843 +    model: Optional[Union[str, LanguageModelLike]] = None,
844 +    subagents: List[SubAgent] = None,
845 +    additional_tools: List[Union[BaseTool, Dict[str, Any]]] = None,
846 +    database_url: Optional[str] = None,
847 +):
848 diff --git a/examples/hacs_developer_agent/graph.py b/examples/hacs_developer_agent/graph.py
849 index 2e84a99..ce15384 100644
850 --- a/examples/hacs_developer_agent/graph.py
851 +++ b/examples/hacs_developer_agent/graph.py
852 @@ -1,212 +1,87 @@
853 -async def admin_operation_router(state: State) -> Dict[str, Any]:
854 -    """Route to appropriate admin operation based on messages."""
855 -    global _current_operation, _operation_in_progress
856 -
857 -    # Get the last message to determine operation type
858 -    messages = state.get("messages", [])
859 -    if not messages:
860 -        return {"last_error": "No operation specified"}
861 -
862 -    last_message = messages[-1].content.lower()
863 -
864 -    # Determine operation type from user message
865 -    if any(word in last_message for word in ["migration", "migrate", "setup", "initialize"]):
866 -        operation_type = "database_migration"
867 -    elif any(word in last_message for word in ["status", "check", "migration status"]):
868 -        operation_type = "migration_status"
869 -    elif any(word in last_message for word in ["schema", "describe", "tables", "structure"]):
870 -        operation_type = "schema_inspection"
871 -    elif any(word in last_message for word in ["create", "record", "resource"]):
872 -        operation_type = "create_resource"
873 -    elif any(word in last_message for word in ["discover", "find", "available"]):
874 -        operation_type = "discover_resources"
875 -    else:
876 -        operation_type = "general_help"
877 -
878 -    # Set private tracking variables
879 -    _current_operation = operation_type
880 -    _operation_in_progress = True
881 -
882 -    return {} # No state changes, just private tracking
883
884 means that the probability distribution is truncated such that as soon as the cumulative probability of the most probable
885 tokens equals or exceeds the top-p value, all other tokens are dropped. In short, lowering temperature decreases the
886 degree to which randomness is introduced in token sampling, while lowering top probability decreases the sample size,
887 making it impossible to sample highly improbable tokens. If agents are interacting frequently with tools, the mitigation
888 of unusual tokens will greatly reduce the likelihood of type errors and tool misuse. Thus, well configured temperature
889 and top probability arguments can improve the reliability of LangChain and LangGraph applications.

```

```

885 Listing 6. Commit 0ad71e1 from ai-resume-agent: Maintain temperature=0.1 for consistent responses; Configure ChatGroq with
886 top_p parameter to reduce creativity
887 diff --git a/app/core/config.py b/app/core/config.py
888 index 3f4f21c..b5a276c 100644
889 --- a/app/core/config.py
890 +++ b/app/core/config.py
891 @@ -2,68 +2,90 @@
892 - GROQ_API_KEY: str
893 - GROQ_MODEL: str = "llama-3.3-70b-versatile" # Modelo actualizado
894 - GROQ_TEMPERATURE: float = 0.7
895 - GROQ_MAX_TOKENS: int = 1024
896 - GROQ\_TIMEOUT: int = 30 # Timeout en segundos (protección anti-DoS)
897 + # Google Gemini API (LLM alternativo)
898 + GEMINI_API_KEY: str = ""
899 + GEMINI_MODEL: str = "gemini-2.5-flash" # Modelo más rápido y menos restrictivo
900 + GEMINI_TEMPERATURE: float = 0.4 # MÁS confianza para sintetizar respuestas STAR
901 + GEMINI_TOP_P: float = 0.7 # Ventana más amplia para construcción de frases
902 + GEMINI_MAX_TOKENS: int = 512 # Espacio suficiente para respuestas detalladas
903
904
905
906
907
908
909
910
911
912
913
914 In the context of a chat bot it may be ideal to introduce a degree of randomness in order to produce more natural
915 language. However, in certain agentic systems, particularly those with tools or specific input and output schemas,
916 developers seek to create a more deterministic system in which the most likely output is typically the desired one.
917 Setting the temperature of a model to zero maximizes the determinism of a model and leads to behavior in which
918 the LLM will always produce the next token with the highest associated probability. With a temperature set to zero,
919 language models will pass arguments to tools with the highest associated probability given the context of the query as
920 well as tool descriptions and type hints. This can enhance the reliability of tool results and potentially mitigate type
921 errors.
922
923
924 Example n: Context trimming for multi-turn threads. When maintaining conversational history in a LangChain or
925 LangGraph application it is common for context to be fed back into the model along with the current query to provide
926 the model with a broader context including previous queries and responses. The act of feeding previous turns back into
927 the model increases the number of tokens being processed by the model and inherently the cost per query. Particularly
928 in the event a conversation between a user and agent takes many turns this will cause the cost of each subsequent
929 query to grow uncontrollably. In extreme cases this may lead to a user exceeding the context window of the LLM. This
930 commit exemplifies a common strategy among developers to trim conversational context to the last n messages. This
931 allows developers to manage the tradeoff between context and cost. It may be that conversational history between the
932 user and agent can enhance the accuracy of future responses at the expense of increased token consumption.
933
934
935 Example n: Prompt Engineering with ChatPromptTemplate. The ChatPromptTemplate object allows developers to
936 organize messages in a multi-turn conversation as a dictionary of strings associated with the correct role. System prompts
937 are correctly identified as such. User generated queries are assigned the role "user" and model generated responses are
938 assigned the role key "assistant". This is advantageous during a multi-turn conversational workflow because it effectively
939 differentiates between each turn and assigns that string to the correct role thus improving model comprehension of
940 context. The prompt template also enables consistent formatting of model responses. ChatPromptTemplate adds support
941 for placeholders allowing for the insertion of context specific data during runtime. The use of a prompt template enables
942 developers to capture context over the course of a multi-turn interaction in order to enhance model awareness.
943
944 Manuscript submitted to ACM

```

```

937 Listing 7. Commit 84365f4 from hikizan-emacs: Implemented message trimming ('trim_messages_window') and memory optimization
938 to manage context window size and prevent overflow.
939 diff --git a/python/emacsagent.py b/python/emacsagent.py
940 index 4b9308f..11ba7be 100644
941 --- a/python/emacsagent.py
942 +++ b/python/emacsagent.py
943 @@ -23,6 +23,8 @@
944     MODEL_NAME = "gemini-2.5-flash"
945     LLM_TEMPERATURE = 0
946     TIMEOUT = 60
947     +MEMORY_WINDOW_SIZE = 10 # Keep last 10 messages in active memory
948     +MAX_EXECUTION_LOG_SIZE = 50 # Keep last 50 execution steps
949     @@ -116,6 +118,100 @@
950     +def trim_messages_window(messages: Sequence[BaseMessage], window_size: int = MEMORY_WINDOW_SIZE) ->
951         Sequence[BaseMessage]:
952     +    """
953     +        Trim messages to keep only the most recent ones within the window size.
954     +        Always preserve the system message if it exists.
955     +    """
956     +    if len(messages) <= window_size:
957         +        return messages
958     +
959     +    # Check if first message is system message
960     +    if messages and isinstance(messages[0], SystemMessage):
961         +        # Keep system message + last (window_size - 1) messages
962         +        return [messages[0]] + list(messages[-(window_size-1):])
963     +    else:
964         +        # Just keep last window_size messages
965         +        return list(messages[-window_size:])
966
967
968
969
970
971
972
973
974
975
976
977
978
979
980
981
982
983
984
985
986
987
988

```

Example n: Initializing models with generic model wrapper for enhanced maintainability. Listing 9 depicts a self-admitted refactoring in which the author simplifies the process of model initialization by using LangChain's generic init_chat_model object. This refactoring allows them to remove three helper functions associated with each of the proprietary providers available in their application. The improved implementation leverages a generic interface while still allowing the flexibility of passing provider-specific arguments via the kwargs field. The removal of hard-coded logic and provider specific functions not only streamlines the function but also expands the number of configurable models to all of those supported by the init_chat_model interface. If a developer wished to add another model option to the application in the previous implementation, they would have to build yet another function and all of the necessary implementation for those unique provider fields. In the new implementation, most mainstream models are available by default.

Example n: Caching model responses for repeat queries. In Listing 10 the author claims to have produced a significant latency reduction by caching cover-letter query responses with a time-to-live of eight hours. In their commit message they explicitly claim to have reduced latency on repeated cover-letter queries from 8 seconds to 3 milliseconds. Although it is difficult to validate this claim without completely cloning the repository and environment, we can logically conclude that caching responses would reduce latency and cost by avoiding an unnecessary external call to a language model. However, this must be done at the expense of increased resources to maintain cached responses for 8 hours at a time. Nonetheless, systems that are prone to receive repeat queries from users and that prioritize quick response time over optimized memory usage will benefit from such an implementation.

```

989 Listing 8. Commit 447a09c from file-classifier: refactor: migrate prompt management to LangChain ChatPromptTemplate
990 diff --git a/src/ai_file_classifier/ai_client.py b/src/ai_file_classifier/ai_client.py
991 index c757756..5a79f45 100644
992 --- a/src/ai_file_classifier/ai_client.py
993 +++ b/src/ai_file_classifier/ai_client.py
994 @@ -133,15 +134,15 @@ def _initialize_llm(
995     system_prompt: str,
996     user_prompt: str,
997     prompt_template: ChatPromptTemplate,
998     prompt_values: dict,
999 new file mode 100644
1000 index 0000000..72a5851
1000 --- /dev/null
1001 +++ b/src/ai_file_classifier/prompt_manager.py
1002 @@ -0,0 +1,73 @@
1003 +def load_file_analysis_prompt() -> ChatPromptTemplate:
1004 +    """
1005 +        Load the file analysis prompt template from text files.
1006 +
1007 +    Returns:
1008 +        ChatPromptTemplate: A LangChain prompt template with system and human messages.
1009 +
1010 +    Raises:
1011 +        FileNotFoundError: If prompt files are missing.
1012 +        IOError: If prompt files cannot be read.
1013 +    """
1014 +    try:
1015 +        # Create ChatPromptTemplate with proper message roles
1016 +        return ChatPromptTemplate.from_messages([
1017 +            ("system", system_prompt),
1018 +            ("human", user_prompt),
1019 +        ])
1020
1021 Example n: Return only state updates from nodes. Per LangGraph documentation nodes are intended to return only
1022 the updated fields of state as a dictionary. These updates are then merged into the existing state by the LangGraph
1023 engine. Depicted in Listing 11 is a refactor in which the author corrected their node function to follow this pattern.
1024 Direct mutations to state can appear to work in some simple graphs that aren't concerned with previous iterations,
1025 however, there are some major risks that can become apparent in complex applications. Firstly, when directly mutating
1026 fields in the state you override any previous values essentially deleting any data from previous nodes or iterations.
1027 This practice may actually function if the program is only ever concerned with the state from the previous step, but
1028 it does not fully encompass the ability of state to maintain data throughout graph execution. Additionally, if a graph
1029 contains branching and parallelism there is a major risk of overriding state updates from parallel processes. If some
1030 state updates are completely overridden and lost, this can hinder graph execution that relies on aspects of state down
1031 the line because now there is missing data.
1032
1033 Example n: Limit API calls to reduce cost and latency. As we defined earlier cost as technical debt in the scope of
1034 this paper refers to the monetary cost associated with large language model queries. We can reduce this cost by
1035 either reducing the number of tokens present within queries or by reducing the number of queries made. Depicted in
1036 Listing 12 is a refactor that elects to limit the number of API calls made to an external provider (OpenAI in this case) by
1037 implementing a conditional edge with simple routing logic. The _should_continue_after_tools function acts as the logic
1038
1039 Manuscript submitted to ACM
1040

```

```

1041 Listing 9. Commit 8c4678d from diversifier: Refactor llm_factory.py to use LangChain's init_chat_model for unified provider
1042 initialization
1043 diff --git a/src/orchestration/llm_factory.py b/src/orchestration/llm_factory.py
1044 index aa13ef8..c84ce2d 100644
1045 --- a/src/orchestration/llm_factory.py
1046 +++ b/src/orchestration/llm_factory.py
1047 @@ -27,153 +28,46 @@ def create_llm_from_config(config: Optional[LLMConfig] = None) -> Any:
1048     # Base parameters common to most providers
1049     params = {
1050         "model_name": config.model_name,
1051         "temperature": config.temperature,
1052         "max_tokens": config.max_tokens,
1053         "timeout": config.timeout,
1054         **config.additional_params,
1055     }
1056     # Add base_url if specified
1057     if config.base_url:
1058         params["base_url"] = config.base_url
1059     logger.info(
1060         f"Creating LLM instance: provider={provider}, model={config.model_name}"
1061     )
1062     if provider == "anthropic":
1063         return _create_anthropic_llm(config, params)
1064     elif provider == "openai":
1065         return _create_openai_llm(config, params)
1066     elif provider == "google":
1067         return _create_google_llm(config, params)
1068     else:
1069         raise ValueError(f"Unsupported LLM provider: {provider}")
1070
1071     # Use init_chat_model with our configuration
1072     from typing import Any
1073
1074     kwargs: dict[str, Any] = {
1075         "temperature": config.temperature,
1076         "max_tokens": config.max_tokens,
1077     }
1078     kwargs.update(config.additional_params)
1079
1080     return init_chat_model(model=model_id, **kwargs)
1081
1082
1083 behind the conditional edge added to the workflow from tools that routes either back to the agent node for further
1084 reasoning or to the synthesize node to complete the workflow. In Listing 12 we observe simple if-else logic that routes
1085 directly to synthesize if the workflow has already successfully executed multiple successful tool calls or iterations. If
1086 neither of these conditions are met then the workflow does return to the agent node and proceeds to call the model for
1087 further reasoning. This function prevents uncontrollably iterating between tools and reasoning which may result in
1088 uncontrolled spending to the API and an increase in latency perceived by the user.
1089
1090
1091 Example n: Initialization Arguments. This particular repository leverages a langchain prebuilt agent for even further
1092 abstraction. As such, the agent executor acts as an object similar to a model when bound with tools. The agent
1093 executor object accepts similar arguments to chat model wrappers such as maximum iterations. In this refactoring the
1094 developer elects to set maximum iterations to 10 where it was not previously. This makes it impossible for the agent to
1095 iterate uncontrollably between tools and the model, repeatedly reasoning about previous tool calls increasing cost and
1096 latency without guardrails. Setting the maximum iterations parameter to an integer depending on the use case of the
1097
1098
1099
1100
1101
1102
1103
1104
1105
1106
1107
1108
1109
1110
1111
1112
1113
1114
1115
1116
1117
1118
1119
1120
1121
1122
1123
1124
1125
1126
1127
1128
1129
1130
1131
1132
1133
1134
1135
1136
1137
1138
1139
1140
1141
1142
1143
1144
1145
1146
1147
1148
1149
1150
1151
1152
1153
1154
1155
1156
1157
1158
1159
1160
1161
1162
1163
1164
1165
1166
1167
1168
1169
1170
1171
1172
1173
1174
1175
1176
1177
1178
1179
1180
1181
1182
1183
1184
1185
1186
1187
1188
1189
1190
1191
1192
1193
1194
1195
1196
1197
1198
1199
1200
1201
1202
1203
1204
1205
1206
1207
1208
1209
1210
1211
1212
1213
1214
1215
1216
1217
1218
1219
1220
1221
1222
1223
1224
1225
1226
1227
1228
1229
1230
1231
1232
1233
1234
1235
1236
1237
1238
1239
1240
1241
1242
1243
1244
1245
1246
1247
1248
1249
1250
1251
1252
1253
1254
1255
1256
1257
1258
1259
1260
1261
1262
1263
1264
1265
1266
1267
1268
1269
1270
1271
1272
1273
1274
1275
1276
1277
1278
1279
1280
1281
1282
1283
1284
1285
1286
1287
1288
1289
1290
1291
1292
1293
1294
1295
1296
1297
1298
1299
1300
1301
1302
1303
1304
1305
1306
1307
1308
1309
1310
1311
1312
1313
1314
1315
1316
1317
1318
1319
1320
1321
1322
1323
1324
1325
1326
1327
1328
1329
1330
1331
1332
1333
1334
1335
1336
1337
1338
1339
1340
1341
1342
1343
1344
1345
1346
1347
1348
1349
1350
1351
1352
1353
1354
1355
1356
1357
1358
1359
1360
1361
1362
1363
1364
1365
1366
1367
1368
1369
1370
1371
1372
1373
1374
1375
1376
1377
1378
1379
1380
1381
1382
1383
1384
1385
1386
1387
1388
1389
1390
1391
1392
1393
1394
1395
1396
1397
1398
1399
1400
1401
1402
1403
1404
1405
1406
1407
1408
1409
1410
1411
1412
1413
1414
1415
1416
1417
1418
1419
1420
1421
1422
1423
1424
1425
1426
1427
1428
1429
1430
1431
1432
1433
1434
1435
1436
1437
1438
1439
1440
1441
1442
1443
1444
1445
1446
1447
1448
1449
1450
1451
1452
1453
1454
1455
1456
1457
1458
1459
1460
1461
1462
1463
1464
1465
1466
1467
1468
1469
1470
1471
1472
1473
1474
1475
1476
1477
1478
1479
1480
1481
1482
1483
1484
1485
1486
1487
1488
1489
1490
1491
1492
1493
1494
1495
1496
1497
1498
1499
1500
1501
1502
1503
1504
1505
1506
1507
1508
1509
1510
1511
1512
1513
1514
1515
1516
1517
1518
1519
1520
1521
1522
1523
1524
1525
1526
1527
1528
1529
1530
1531
1532
1533
1534
1535
1536
1537
1538
1539
1540
1541
1542
1543
1544
1545
1546
1547
1548
1549
1550
1551
1552
1553
1554
1555
1556
1557
1558
1559
1560
1561
1562
1563
1564
1565
1566
1567
1568
1569
1570
1571
1572
1573
1574
1575
1576
1577
1578
1579
1580
1581
1582
1583
1584
1585
1586
1587
1588
1589
1590
1591
1592
1593
1594
1595
1596
1597
1598
1599
1600
1601
1602
1603
1604
1605
1606
1607
1608
1609
1610
1611
1612
1613
1614
1615
1616
1617
1618
1619
1620
1621
1622
1623
1624
1625
1626
1627
1628
1629
1630
1631
1632
1633
1634
1635
1636
1637
1638
1639
1640
1641
1642
1643
1644
1645
1646
1647
1648
1649
1650
1651
1652
1653
1654
1655
1656
1657
1658
1659
1660
1661
1662
1663
1664
1665
1666
1667
1668
1669
1670
1671
1672
1673
1674
1675
1676
1677
1678
1679
1680
1681
1682
1683
1684
1685
1686
1687
1688
1689
1690
1691
1692
1693
1694
1695
1696
1697
1698
1699
1700
1701
1702
1703
1704
1705
1706
1707
1708
1709
1710
1711
1712
1713
1714
1715
1716
1717
1718
1719
1720
1721
1722
1723
1724
1725
1726
1727
1728
1729
1730
1731
1732
1733
1734
1735
1736
1737
1738
1739
1740
1741
1742
1743
1744
1745
1746
1747
1748
1749
1750
1751
1752
1753
1754
1755
1756
1757
1758
1759
1760
1761
1762
1763
1764
1765
1766
1767
1768
1769
1770
1771
1772
1773
1774
1775
1776
1777
1778
1779
1780
1781
1782
1783
1784
1785
1786
1787
1788
1789
1790
1791
1792
1793
1794
1795
1796
1797
1798
1799
1800
1801
1802
1803
1804
1805
1806
1807
1808
1809
1810
1811
1812
1813
1814
1815
1816
1817
1818
1819
1820
1821
1822
1823
1824
1825
1826
1827
1828
1829
1830
1831
1832
1833
1834
1835
1836
1837
1838
1839
1840
1841
1842
1843
1844
1845
1846
1847
1848
1849
1850
1851
1852
1853
1854
1855
1856
1857
1858
1859
1860
1861
1862
1863
1864
1865
1866
1867
1868
1869
1870
1871
1872
1873
1874
1875
1876
1877
1878
1879
1880
1881
1882
1883
1884
1885
1886
1887
1888
1889
1889
1890
1891
1892
1893
1894
1895
1896
1897
1898
1899
1900
1901
1902
1903
1904
1905
1906
1907
1908
1909
1909
1910
1911
1912
1913
1914
1915
1916
1917
1918
1919
1920
1921
1922
1923
1924
1925
1926
1927
1928
1929
1929
1930
1931
1932
1933
1934
1935
1936
1937
1938
1939
1939
1940
1941
1942
1943
1944
1945
1946
1947
1948
1949
1949
1950
1951
1952
1953
1954
1955
1956
1957
1958
1959
1959
1960
1961
1962
1963
1964
1965
1966
1967
1968
1969
1969
1970
1971
1972
1973
1974
1975
1976
1977
1978
1979
1979
1980
1981
1982
1983
1984
1985
1986
1987
1988
1989
1989
1990
1991
1992
1993
1994
1995
1996
1997
1998
1999
1999
2000
2001
2002
2003
2004
2005
2006
2007
2008
2009
2009
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2099
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2099
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2099
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2099
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2099
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2099
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2099
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2099
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2099
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2099
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2099
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2089
2090
2091
2092
2093
2094
2095
2096
209
```

```

1093      Listing 10. Commit 8108951 from aruizca-resume: Cover letter caching reduces 8s operations to 3ms (cache hits)
1094 diff --git a/packages/core/src/main/cover-letter/infrastructure/langchain/CoverLetterPromptRunner.ts b
1095   /packages/core/src/main/cover-letter/infrastructure/langchain/CoverLetterPromptRunner.ts
1096 index 5866671d..18ca0db 100644
1097 --- a/packages/core/src/main/cover-letter/infrastructure/langchain/CoverLetterPromptRunner.ts
1098 +++ b/packages/core/src/main/cover-letter/infrastructure/langchain/CoverLetterPromptRunner.ts
1099 +   cacheConfig: {
1100 +     ttl: 8 * 60 * 60 * 1000 // 8 hours
1101 +   },
1102 diff --git a/packages/core/src/main/shared/infrastructure/langchain/LangchainPromptRunner.ts b/
1103   packages/core/src/main/shared/infrastructure/langchain/LangchainPromptRunner.ts
1104 new file mode 100644
1105 index 0000000..74d22a5
1106 --- /dev/null
1107 +++ b/packages/core/src/main/shared/infrastructure/langchain/LangchainPromptRunner.ts
1108 @@ -0,0 +1,153 @@
1109 +   // Check cache first
1110 +   const cachedResponse = await this.cache.get(input, promptTemplateString, forceRefresh);
1111 +   if (cachedResponse && !forceRefresh) {
1112 +     return this.config.outputTransformer(cachedResponse);
1113 +   }
1114 +   // Execute the LLM operation with performance monitoring
1115 +   const result = await performanceMonitor.trackOperation(
1116 +     operationName,
1117 +     async () => {
1118 +       // Create and execute chain
1119 +       const chain = await this.createChain();
1120 +       return await chain.invoke(promptVariables);
1121 +     },
1122 +     { logToConsole: true }
1123 +   );
1124 +   // Cache the response
1125 +   await this.cache.set(input, promptTemplateString, result);
1126
1127      Listing 11. Commit c373419 from JAA: Updated Node Return Values - All nodes now return only state updates (dictionaries) instead
1128 of mutating and returning full state objects
1129 diff --git a/l1m/l1m.py b/l1m/l1m.py
1130 index 4a6fea4..f2bf4d9 100644
1131 --- a/l1m/l1m.py
1132 +++ b/l1m/l1m.py
1133 @@ -195,12 +214,14 @@ def _router_node(self, state: AgentState) -> AgentState:
1134     else:
1135         func_name = RouterFunction.HUMAN_TEXT
1136         query = ""
1137         state["router_decision"] = func_name
1138         state["router_query"] = query
1139         state["query_keywords"] = keywords
1140
1141         return state
1142     # Return only updates - LangGraph will merge into state
1143     return {
1144         "router_decision": func_name,
1145         "router_query": query,
1146         "query_keywords": keywords
1147     }
1148
1149      application is a near effortless action that prevents behaviors that can negatively impact the cost and latency of an
1150      agentic application.
1151
1152 Manuscript submitted to ACM

```

```
1145 Listing 12. Commit 0760805 from cybershield: Reduce OpenAI API calls by 75% (from 4-8+ to 1-2 calls per analysis); Implement smart
1146 synthesis routing after sufficient tool results
1147 diff --git a/workflows/react_workflow.py b/workflows/react_workflow.py
1148 index d22f0e1..915e359 100644
1149 --- a/workflows/react_workflow.py
1150 +++ b/workflows/react_workflow.py
1151 @@ -73,17 +83,24 @@
1152     # Add edges - optimized to reduce OpenAI API calls
1153     workflow.set_entry_point("agent")
1154     workflow.add_conditional_edges(
1155         "agent",
1156         self._should_continue,
1157         {
1158             "continue": "tools",
1159             "end": "synthesize"
1160         }
1161     )
1162     workflow.add_conditional_edges(
1163         "tools",
1164         self._should_continue_after_tools,
1165         {
1166             "agent": "agent",      # Only go back to agent if more reasoning needed
1167             "synthesize": "synthesize" # Otherwise synthesize directly
1168         }
1169     )
1170     workflow.add_edge("tools", "agent")
1171 @@ -370,21 +387,49 @@
1172     def _should_continue_after_tools(self, state: CyberShieldState) -> str:
1173         """Decide whether to continue reasoning or synthesize after tool execution"""
1174         iteration = state.get("iteration_count", 0)
1175         scratchpad = state.get("agent_scratchpad", "")
1176         # Count successful tool executions
1177         successful_observations = scratchpad.count("Observation:") - scratchpad.count('error')
1178         # If we have multiple successful tool results, go straight to synthesis
1179         if successful_observations >= 2 or iteration >= 1:
1180             logger.info("Moving to synthesis after tools",
1181                         successful_observations=successful_observations,
1182                         iteration=iteration,
1183                         reason="sufficient_data")
1184         return "synthesize"
1185     elif iteration > 5: # Hard limit to prevent loops
1186         logger.info("Forcing synthesis due to iteration limit", iteration=iteration)
1187         return "synthesize"
1188     else:
1189         return "continue"
1190     logger.info("Continuing to agent for more reasoning", iteration=iteration)
1191     return "agent"
```

Example n: Tool Correction with descriptive naming & docstring. As described in Table 3, the tool decorator from LangChain is designed to use the Docstring of the tool function as a description for the model. This allows the model to reason about the tool's functionality and determine if it's use-cases align with the query. As a result, developers are encouraged to make a functionally descriptive docstring with the elements "Args" and "Returns" included. In the event that the model does elect to execute this tool, it will now be informed of the arguments that it needs to pass and what is expected to be returned. This greatly enhances the reliability of model actions because it reduces the likelihood

```

1197          Listing 13. Commit e50b85e from ai-cv-agent: Refactor LangChain CV Agent
1198
1199      diff --git a/agent/langchain_cv_agent.py b/agent/langchain_cv_agent.py
1200      index 20dc2b4..cb0eddd 100644
1201      --- a/agent/langchain_cv_agent.py
1202      +++ b/agent/langchain_cv_agent.py
1203      @@ -235,32 +201,51 @@
1204      -         return AgentExecutor(agent=agent, tools=self.tools, verbose=True)
1205      +     return AgentExecutor(
1206      +         agent=agent,
1207      +         tools=self.tools,
1208      +         verbose=True,
1209      +         max_iterations=10,
1210      +         handle_parsing_errors=True,
1211      +         return_intermediate_steps=True # Enable to see artifact passing
1212      +
1213
1214
1215          Listing 14. Commit e50b85e from ai-cv-agent: Refactor LangChain CV Agent
1216
1217      diff --git a/agent/langchain_cv_agent.py b/agent/langchain_cv_agent.py
1218      index 20dc2b4..cb0eddd 100644
1219      --- a/agent/langchain_cv_agent.py
1220      +++ b/agent/langchain_cv_agent.py
1221      @@ -51,179 +57,139 @@
1222      -     @tool
1223      -     def analyze_job_posting(job_url: str) -> str:
1224      -         """Analyze a job posting and extract key requirements, skills, and keywords"""
1225      -         @tool(response_format="content_and_artifact")
1226      -     def analyze_job_with_llm(job_description: str) -> tuple[str, dict]:
1227      -         """
1228      -             Analyze job posting to extract key requirements and keywords.
1229      -             Args:
1230      -                 job_description: The job description text
1231      -             Returns:
1232      -                 Tuple of (analysis summary, structured analysis dict as artifact)
1233      -             """
1234
1235
1236
1237
1238
1239
1240
1241
1242
1243
1244
1245
1246
1247
1248
1249
1250
1251
1252
1253
1254
1255
1256
1257
1258
1259
1260
1261
1262
1263
1264
1265
1266
1267
1268
1269
1270
1271
1272
1273
1274
1275
1276
1277
1278
1279
1280
1281
1282
1283
1284
1285
1286
1287
1288
1289
1290
1291
1292
1293
1294
1295
1296
1297
1298
1299
1300
1301
1302
1303
1304
1305
1306
1307
1308
1309
1310
1311
1312
1313
1314
1315
1316
1317
1318
1319
1320
1321
1322
1323
1324
1325
1326
1327
1328
1329
1330
1331
1332
1333
1334
1335
1336
1337
1338
1339
1340
1341
1342
1343
1344
1345
1346
1347
1348
1349
1350
1351
1352
1353
1354
1355
1356
1357
1358
1359
1360
1361
1362
1363
1364
1365
1366
1367
1368
1369
1370
1371
1372
1373
1374
1375
1376
1377
1378
1379
1380
1381
1382
1383
1384
1385
1386
1387
1388
1389
1390
1391
1392
1393
1394
1395
1396
1397
1398
1399
1400
1401
1402
1403
1404
1405
1406
1407
1408
1409
1410
1411
1412
1413
1414
1415
1416
1417
1418
1419
1420
1421
1422
1423
1424
1425
1426
1427
1428
1429
1430
1431
1432
1433
1434
1435
1436
1437
1438
1439
1440
1441
1442
1443
1444
1445
1446
1447
1448
1449
1450
1451
1452
1453
1454
1455
1456
1457
1458
1459
1460
1461
1462
1463
1464
1465
1466
1467
1468
1469
1470
1471
1472
1473
1474
1475
1476
1477
1478
1479
1480
1481
1482
1483
1484
1485
1486
1487
1488
1489
1490
1491
1492
1493
1494
1495
1496
1497
1498
1499
1500
1501
1502
1503
1504
1505
1506
1507
1508
1509
1510
1511
1512
1513
1514
1515
1516
1517
1518
1519
1520
1521
1522
1523
1524
1525
1526
1527
1528
1529
1530
1531
1532
1533
1534
1535
1536
1537
1538
1539
1540
1541
1542
1543
1544
1545
1546
1547
1548
1549
1550
1551
1552
1553
1554
1555
1556
1557
1558
1559
1560
1561
1562
1563
1564
1565
1566
1567
1568
1569
1570
1571
1572
1573
1574
1575
1576
1577
1578
1579
1580
1581
1582
1583
1584
1585
1586
1587
1588
1589
1590
1591
1592
1593
1594
1595
1596
1597
1598
1599
1600
1601
1602
1603
1604
1605
1606
1607
1608
1609
1610
1611
1612
1613
1614
1615
1616
1617
1618
1619
1620
1621
1622
1623
1624
1625
1626
1627
1628
1629
1630
1631
1632
1633
1634
1635
1636
1637
1638
1639
1640
1641
1642
1643
1644
1645
1646
1647
1648
1649
1650
1651
1652
1653
1654
1655
1656
1657
1658
1659
1660
1661
1662
1663
1664
1665
1666
1667
1668
1669
1670
1671
1672
1673
1674
1675
1676
1677
1678
1679
1680
1681
1682
1683
1684
1685
1686
1687
1688
1689
1690
1691
1692
1693
1694
1695
1696
1697
1698
1699
1700
1701
1702
1703
1704
1705
1706
1707
1708
1709
1710
1711
1712
1713
1714
1715
1716
1717
1718
1719
1720
1721
1722
1723
1724
1725
1726
1727
1728
1729
1730
1731
1732
1733
1734
1735
1736
1737
1738
1739
1740
1741
1742
1743
1744
1745
1746
1747
1748
1749
1750
1751
1752
1753
1754
1755
1756
1757
1758
1759
1760
1761
1762
1763
1764
1765
1766
1767
1768
1769
1770
1771
1772
1773
1774
1775
1776
1777
1778
1779
1780
1781
1782
1783
1784
1785
1786
1787
1788
1789
1790
1791
1792
1793
1794
1795
1796
1797
1798
1799
1800
1801
1802
1803
1804
1805
1806
1807
1808
1809
1810
1811
1812
1813
1814
1815
1816
1817
1818
1819
1820
1821
1822
1823
1824
1825
1826
1827
1828
1829
1830
1831
1832
1833
1834
1835
1836
1837
1838
1839
1840
1841
1842
1843
1844
1845
1846
1847
1848
1849
1850
1851
1852
1853
1854
1855
1856
1857
1858
1859
1860
1861
1862
1863
1864
1865
1866
1867
1868
1869
1870
1871
1872
1873
1874
1875
1876
1877
1878
1879
1880
1881
1882
1883
1884
1885
1886
1887
1888
1889
1890
1891
1892
1893
1894
1895
1896
1897
1898
1899
1900
1901
1902
1903
1904
1905
1906
1907
1908
1909
1910
1911
1912
1913
1914
1915
1916
1917
1918
1919
1920
1921
1922
1923
1924
1925
1926
1927
1928
1929
1930
1931
1932
1933
1934
1935
1936
1937
1938
1939
1940
1941
1942
1943
1944
1945
1946
1947
1948
1949
1950
1951
1952
1953
1954
1955
1956
1957
1958
1959
1960
1961
1962
1963
1964
1965
1966
1967
1968
1969
1970
1971
1972
1973
1974
1975
1976
1977
1978
1979
1980
1981
1982
1983
1984
1985
1986
1987
1988
1989
1990
1991
1992
1993
1994
1995
1996
1997
1998
1999
2000
2001
2002
2003
2004
2005
2006
2007
2008
2009
2010
2011
2012
2013
2014
2015
2016
2017
2018
2019
2020
2021
2022
2023
2024
2025
2026
2027
2028
2029
2030
2031
2032
2033
2034
2035
2036
2037
2038
2039
2040
2041
2042
2043
2044
2045
2046
2047
2048
2049
2050
2051
2052
2053
2054
2055
2056
2057
2058
2059
2060
2061
2062
2063
2064
2065
2066
2067
2068
2069
2070
2071
2072
2073
2074
2075
2076
2077
2078
2079
2080
2081
2082
2083
2084
2085
2086
2087
2088
2089
2090
2091
2092
2093
2094
2095
2096
2097
2098
2099
2100
2101
2102
2103
2104
2105
2106
2107
2108
2109
2110
2111
2112
2113
2114
2115
2116
2117
2118
2119
2120
2121
2122
2123
2124
2125
2126
2127
2128
2129
2130
2131
2132
2133
2134
2135
2136
2137
2138
2139
2140
2141
2142
2143
2144
2145
2146
2147
2148
2149
2150
2151
2152
2153
2154
2155
2156
2157
2158
2159
2160
2161
2162
2163
2164
2165
2166
2167
2168
2169
2170
2171
2172
2173
2174
2175
2176
2177
2178
2179
2180
2181
2182
2183
2184
2185
2186
2187
2188
2189
2190
2191
2192
2193
2194
2195
2196
2197
2198
2199
2200
2201
2202
2203
2204
2205
2206
2207
2208
2209
2210
2211
2212
2213
2214
2215
2216
2217
2218
2219
2220
2221
2222
2223
2224
2225
2226
2227
2228
2229
2230
2231
2232
2233
2234
2235
2236
2237
2238
2239
2240
2241
2242
2243
2244
2245
2246
2247
2248
2249
2250
2251
2252
2253
2254
2255
2256
2257
2258
2259
2260
2261
2262
2263
2264
2265
2266
2267
2268
2269
2270
2271
2272
2273
2274
2275
2276
2277
2278
2279
2280
2281
2282
2283
2284
2285
2286
2287
2288
2289
2290
2291
2292
2293
2294
2295
2296
2297
2298
2299
2300
2301
2302
2303
2304
2305
2306
2307
2308
2309
2310
2311
2312
2313
2314
2315
2316
2317
2318
2319
2320
2321
2322
2323
2324
2325
2326
2327
2328
2329
2330
2331
2332
2333
2334
2335
2336
2337
2338
2339
2340
2341
2342
2343
2344
2345
2346
2347
2348
2349
2350
2351
2352
2353
2354
2355
2356
2357
2358
2359
2360
2361
2362
2363
2364
2365
2366
2367
2368
2369
2370
2371
2372
2373
2374
2375
2376
2377
2378
2379
2380
2381
2382
2383
2384
2385
2386
2387
2388
2389
2390
2391
2392
2393
2394
2395
2396
2397
2398
2399
2400
2401
2402
2403
2404
2405
2406
2407
2408
2409
2410
2411
2412
2413
2414
2415
2416
2417
2418
2419
2420
2421
2422
2423
2424
2425
2426
2427
2428
2429
2430
2431
2432
2433
2434
2435
2436
2437
2438
2439
2440
2441
2442
2443
2444
2445
2446
2447
2448
2449
2450
2451
2452
2453
2454
2455
2456
2457
2458
2459
2460
2461
2462
2463
2464
2465
2466
2467
2468
2469
2470
2471
2472
2473
2474
2475
2476
2477
2478
2479
2480
2481
2482
2483
2484
2485
2486
2487
2488
2489
2490
2491
2492
2493
2494
2495
2496
2497
2498
2499
2500
2501
2502
2503
2504
2505
2506
2507
2508
2509
2510
2511
2512
2513
2514
2515
2516
2517
2518
2519
2520
2521
2522
2523
2524
2525
2526
2527
2528
2529
2530
2531
2532
2533
2534
2535
2536
2537
2538
2539
2540
2541
2542
2543
2544
2545
2546
2547
2548
2549
2550
2551
2552
2553
2554
2555
2556
2557
2558
2559
2560
2561
2562
2563
2564
2565
2566
2567
2568
2569
2570
2571
2572
2573
2574
2575
2576
2577
2578
2579
2580
2581
2582
2583
2584
2585
2586
2587
2588
2589
2590
2591
2592
2593
2594
2595
2596
2597
2598
2599
2600
2601
2602
2603
2604
2605
2606
2607
2608
2609
2610
2611
2612
2613
2614
2615
2616
2617
2618
2619
2620
2621
2622
2623
2624
2625
2626
2627
2628
2629
2630
2631
2632
2633
2634
2635
2636
2637
2638
2639
2640
2641
2642
2643
2644
2645
2646
2647
2648
2649
2650
2651
2652
2653
2654
2655
2656
2657
2658
2659
2660
2661
2662
2663
2664
2665
2666
2667
2668
2669
2670
2671
2672
2673
2674
2675
2676
2677
2678
2679
2680
2681
2682
2683
2684
2685
2686
2687
2688
2689
2690
2691
2692
2693
2694
2695
2696
2697
2698
2699
2700
2701
2702
2703
2704
2705
2706
2707
2708
2709
2710
2711
2712
2713
2714
2715
2716
2717
2718
2719
2720
2721
2722
2723
2724
2725
2726
2727
2728
2729
2730
2731
2732
2733
2734
2735
2736
2737
2738
2739
2740
2741
2742
2743
2744
2745
2746
2747
2748
2749
2750
2751
2752
2753
2754
2755
2756
2757
2758
2759
2760
2761
2762
2763
2764
2765
2766
2767
2768
2769
2770
2771
2772
2773
2774
2775
2776
2777
2778
2779
2780
2781
2782
2783
2784
2785
2786
2787
2788
2789
2790
2791
2792
2793
2794
2795
2796
2797
2798
2799
2800
2801
2802
2803
2804
2805
2806
2807
2808
2809
2810
2811
2812
2813
2814
2815
2816
2817
2818
2819
2820
2821
2822
2823
2824
2825
2826
2827
2828
2829
2830
2831
2832
2833
2834
2835
2836
2837
2838
2839
2840
2841
2842
2843
2844
2845
2846
2847
2848
2849
2850
2851
2852
2853
2854
2855
2856
2857
2858
2859
2860
2861
2862
2863
2864
2865
2866
2867
2868
2869
2870
2871
2872
2873
2874
2875
2876
2877
2878
2879
2880
2881
2882
2883
2884
2885
2886
2887
2888
2889
2890
2891
2892
2893
2894
2895
2896
2897
2898
2899
2900
2901
2902
2903
2904
2905
2906
2907
2908
2909
2910
2911
2912
2913
2914
2915
2916
2917
2918
2919
2920
2921
2922
2923
2924
2925
2926
2927
2928
2929
2930
2931
2932
2933
2934
2935
2936
2937
2938
2939
2940
2941
2942
2943
2944
2945
2946
2947
2948
2949
2950
2951
2952
2953
2954
2955
2956
2957
2958
2959
2960
2961
2962
2963
2964
2965
2966
2967
2968
2969
2970
2971
2972
2973
2974
2975
2976
2977
2978
2979
2980
2981
2982
2983
2984
2985
2986
2987
2988
2989
2990
2991
2992
2993
2994
2995
2996
2997
2998
2999
2999

```

```

1249
1250 Listing 15. Commit b44b80a from medabot: Added maxTokens: 1000 to limit response length
1251 diff --git a/app/core/leafletProcessor.ts b/app/core/leafletProcessor.ts
1252 index c42a33e..ec51887 100644
1253 --- a/app/core/leafletProcessor.ts
1254 +++ b/app/core/leafletProcessor.ts
1255 @@ -50,61 +78,94 @@
1256 - // Create chat model
1257 + // Use gpt-4o-mini - better quality, faster, and cheaper than gpt-4.1-nano
1258 + // gpt-4o-mini has 128k context window and better instruction following
1259 const llm = new ChatOpenAI({
1260 - modelName: "gpt-4.1-nano",
1261 - temperature: 0,
1262 + modelName: "gpt-4o-mini", // Updated from gpt-4.1-nano
1263 + temperature: 0.1, // Slightly increased for more natural language, but still mostly
1264 deterministic
1265 openAIApiKey: process.env.OPENAI_API_KEY,
1266 + maxTokens: 1000, // Limit response length
1267
1268
1269 Listing 16. Commit e7cd19c from Agente-de-IA-usando-Next-y-Langchain: Introduce a trimmer to manage conversation history by
1270 trimming messages based on token count
1271 diff --git a/lib/langgraph.ts b/lib/langgraph.ts
1272 index 74df180..e076539 100644
1273 --- a/lib/langgraph.ts
1274 +++ b/lib/langgraph.ts
1275 @@ -1,12 +1,20 @@
1276 +// Create a trimmer to trim the messages
1277 +const trimmer = trimMessages({
1278 + maxTokens: 10,
1279 + strategy: "last",
1280 + tokenCounter: (msgs) => msgs.length,
1281 + includeSystem: true,
1282 + allowPartial: false,
1283 + startOn: "human",
1284 +});
1285
1286 + // Trim the messages to manage conversation history
1287 + const trimmedMessages = await trimmer.invoke(state.messages);
1288 + // Format the prompt with the current messages
1289 + const prompt = await promptTemplate.invoke({ messages: trimmedMessages });
1290 + const response = await model.invoke(prompt);
1291 + return { messages: [response] };
1292
1293
1294
1295
1296
1297
1298
1299
1300
1301 Example n: Model routing based on perceived query difficulty. Our study has observed repeatedly that repositories in
1302 the open source community will often route to different large language models depending upon the user query. This
1303 can be done with one of two motivations. Firstly, some applications may route to specific models depending upon the
1304 task implicitly stated in a query to take advantage of the strengths of particular models. Secondly, a developer may
1305 intend to extract the difficulty of the query and route to a specific model based on that perceived difficulty; this is the
1306 scenario that is depicted in Listing 17. In this particular project the commit author has attempted to extract the difficulty
1307 of the query based on user input and route to a smaller, less intensive model if the agentic strength is perceived to be
1308 low or very low. In this particular implementation both models (gpt-4o and gpt-4o-mini) are provided by OpenAI, but
1309 the mini model allows for both a cost and latency reduction when queries are simple and extensive reasoning is not
1310 required. This is an optimization that allows for improved cost and performance as technical debt while also allowing
1311 the user the opportunity to use a more robust model for reliable results with more complex tasks.
```

```

1301 Listing 17. Commit 231d4f2 from bt-servant-engine: honor the effective strength by switching to 'gpt-4o-mini' when the level is low,
1302 and introduce the new 'set-agentic-strength' intent that persists user requests
1303 diff --git a/brain.py b/brain.py
1304 index 2f6cd08..9f3868a 100644
1305 --- a/brain.py
1306 +++ b/brain.py
1307 @@ -719,6 +735,68 @@
1308 +def _model_for_agentic_strength(
1309 +    agentic_strength: str,
1310 +    *,
1311 +    allow_low: bool,
1312 +    allow_very_low: bool,
1313 +) -> str:
1314     """Return GPT model name based on strength and allowed downgrades."""
1315     allowed: set[str] = set()
1316     if allow_low:
1317         allowed.add("low")
1318     if allow_very_low:
1319         allowed.add("very_low")
1320     return "gpt-4o-mini" if agentic_strength in allowed else "gpt-4o"
1321 @@ -1271,7 +1674,7 @@
1322     completion = open_ai_client.chat.completions.create(
1323         model="gpt-4o",
1324         model=model_name,
1325         @@@ -1356,7 +1764,8 @@
1326         query_language = detect_language(query)
1327         agentic_strength = _resolve_agentic_strength(s)
1328         query_language = detect_language(query, agentic_strength=agentic_strength)
1329
1330
1331 Listing 18. Commit 818d728 from agent_inbox: Handle both string and list content (multimodal messages); Type-safe content
1332 extraction for AIMessage and ToolMessage
1333 diff --git a/src/graph.py b/src/graph.py
1334 index a4eb0ce..aad0382 100644
1335 --- a/src/graph.py
1336 +++ b/src/graph.py
1337 @@ -676,15 +676,21 @@
1338 -     if isinstance(last_msg, AIMessage) and (
1339 -         "recap" in last_msg.content.lower() or "summary" in last_msg.content.lower()
1340 -     ):
1341 -         return {"messages": messages}
1342 -     # Handle both string and list content (multimodal messages with tool calls)
1343 -     if isinstance(last_msg, AIMessage):
1344 -         content = last_msg.content
1345 -         if isinstance(content, str) and (
1346 -             "recap" in content.lower() or "summary" in content.lower()
1347 -         ):
1348 -             return {"messages": messages}
1349 -         source_text = ""
1350 -         for m in reversed(messages):
1351 -             if isinstance(m, AIMessage) or isinstance(m, ToolMessage):
1352 -                 source_text = (m.content or "").strip()
1353 -                 # Type-safe content extraction for multimodal messages
1354 -                 content = m.content or ""
1355 -                 source_text = content if isinstance(content, str) else str(content)
1356 -                 source_text = source_text.strip()
1357
1358
1359 Example n: Improved message handling. The LangChain handles messages as either HumanMessage, AIMessage, or
1360 ToolMessage objects, all of which extend the BaseMessage class. Each message type contains not only the contents of
1361 Manuscript submitted to ACM

```

1353 a query or response but also the necessary metadata associated with their execution. Corrections made to message
 1354 handling is one of the most frequently observed patterns in LangChain refactorings. The refactor depicted in Listing 18
 1355 alters the post-model hook to handle messages more resiliently regardless of the message type. The refactored method
 1356 allows for graceful handling of tool calls in the event that they are returned while still maintaining the ability to extract
 1357 the contents of an AIMessage object. The improved implementation avoids type-errors in the event that tool calls or
 1358 other metadata are contained within the message object.
 1359

1361 5 Citations and Bibliographies

1363 The use of Bib_TE_X for the preparation and formatting of one's references is strongly recommended. Authors' names
 1364 should be complete – use full first names ("Donald E. Knuth") not initials ("D. E. Knuth") – and the salient identifying
 1365 features of a reference should be included: title, year, volume, number, pages, article DOI, etc.

1366 The bibliography is included in your source document with these two commands, placed just before the \end{document}
 1367 command:

```
1369 \bibliographystyle{ACM-Reference-Format}
1370 \bibliography{bibfile}
```

1372 where "bibfile" is the name, without the ".bib" suffix, of the Bib_TE_X file.

1373 Citations and references are numbered by default. A small number of ACM publications have citations and references
 1374 formatted in the "author year" style; for these exceptions, please include this command in the **preamble** (before the
 1375 command "\begin{document}") of your L_AT_EX source:

```
1377 \citetstyle{acmauthoryear}
```

1379 Some examples. A paginated journal article [2], an enumerated journal article [11], a reference to an entire issue [10],
 1380 a monograph (whole book) [24], a monograph/whole book in a series (see 2a in spec. document) [18], a divisible-book
 1381 such as an anthology or compilation [13] followed by the same example, however we only output the series if the
 1382 volume number is given [14] (so Editor00a's series should NOT be present since it has no vol. no.), a chapter in a divisible
 1383 book [36], a chapter in a divisible book in a series [12], a multi-volume work as book [23], a couple of articles in a
 1384 proceedings (of a conference, symposium, workshop for example) (paginated proceedings article) [3, 16], a proceedings
 1385 article with all possible elements [35], an example of an enumerated proceedings article [15], an informally published
 1386 work [17], a couple of preprints [6, 8], a doctoral dissertation [9], a master's thesis: [4], an online document / world wide
 1387 web resource [1, 28, 37], a video game (Case 1) [27] and (Case 2) [26] and [25] and (Case 3) a patent [34], work accepted
 1388 for publication [31], 'YYYYb'-test for prolific author [32] and [33]. Other cites might contain 'duplicate' DOI and URLs
 1389 (some SIAM articles) [22]. Boris / Barbara Beeton: multi-volume works as books [20] and [19]. A presentation [30]. An
 1390 article under review [7]. A couple of citations with DOIs: [21, 22]. Online citations: [37–39]. Artifacts: [29] and [5].
 1391

1395 6 Acknowledgments

1396 Identification of funding sources and other support, and thanks to individuals and groups that assisted in the research
 1397 and the preparation of the work should be included in an acknowledgment section, which is placed just before the
 1398 reference section in your document.

1400 This section has a special environment:

```
1401 \begin{acks}
1402 ...
1403 ...
```

1405 \end{acks}

1406

1407 so that the information contained therein can be more easily collected during the article metadata extraction phase, and
1408 to ensure consistency in the spelling of the section heading.

1409

1410 Authors should not prepare this section as a numbered or unnumbered \section; please use the “acks” environment.

1411

1412 7 Appendices

1413

1414 If your work needs an appendix, add it before the “\end{document}” command at the conclusion of your source
1415 document.

1416

1417 Start the appendix with the “appendix” command:

1418

1419 \appendix

1420

1421 and note that in the appendix, sections are lettered, not numbered. This document has two appendices, demonstrating
1422 the section and subsection identification method.

1423

1424 8 SIGCHI Extended Abstracts

1425

1426 The “sigchi-a” template style (available only in L^AT_EX and not in Word) produces a landscape-orientation formatted
1427 article, with a wide left margin. Three environments are available for use with the “sigchi-a” template style, and
1428 produce formatted output in the margin:

1429

1430 **sidebar:** Place formatted text in the margin.

1431

1432 **marginfigure:** Place a figure in the margin.

1433

1434 **maintable:** Place a table in the margin.

1435

1436 Acknowledgments

1437

1438 To Robert, for the bagels and explaining CMYK and color spaces.

1439

1440 References

1441

- [1] Rafał Ablamowicz and Bertfried Fauser. 2007. *CLIFFORD: a Maple 11 Package for Clifford Algebra Computations, version 11*. Retrieved February 28, 2008 from <http://math.tntech.edu/rafal/cliff11/index.html>
- [2] Patricia S. Abril and Robert Plant. 2007. The patent holder’s dilemma: Buy, sell, or troll? *Commun. ACM* 50, 1 (Jan. 2007), 36–44. doi:10.1145/1188913.1188915
- [3] Sten Andler. 1979. Predicate Path expressions. In *Proceedings of the 6th ACM SIGACT-SIGPLAN symposium on Principles of Programming Languages (POPL ’79)*. ACM Press, New York, NY, 226–236. doi:10.1145/567752.567774
- [4] David A. Anisi. 2003. *Optimal Motion Control of a Ground Vehicle*. Master’s thesis. Royal Institute of Technology (KTH), Stockholm, Sweden.
- [5] Sam Anzaroot and Andrew McCallum. 2013. *UMass Citation Field Extraction Dataset*. Retrieved May 27, 2019 from <http://www.iesl.cs.umass.edu/data/data-umasscitationfield>
- [6] Sam Anzaroot, Alexandre Passos, David Belanger, and Andrew McCallum. 2014. *Learning Soft Linear Constraints with Application to Citation Field Extraction*. arXiv:1403.1349 doi:10.48550/arXiv.1403.1349
- [7] R. Baggett, M. Simecek, C. Chambellan, K. Tsui, and M. Fraune. 2025. Fluidity in the Phased Framework of Technology Acceptance: Case Study to Gain a Holistic Understanding of (Older Adult) Participant Advancement Through Acceptance Phases with Mobile Telepresence Robots. *Robotics Aut. Systems*. Manuscript submitted for review.
- [8] Lutz Bornmann, K. Brad Wray, and Robin Haunschild. 2019. *Citation concept analysis (CCA)—A new form of citation analysis revealing the usefulness of concepts for other researchers illustrated by two exemplary case studies including classic books by Thomas S. Kuhn and Karl R. Popper*. arXiv:1905.12410 [cs.DL]
- [9] Kenneth L. Clarkson. 1985. *Algorithms for Closest-Point Problems (Computational Geometry)*. Ph.D. Dissertation. Stanford University, Palo Alto, CA. UMI Order Number: AAT 8506171.
- [10] Jacques Cohen (Ed.). 1996. Special issue: Digital Libraries. *Commun. ACM* 39, 11 (Nov. 1996).

1450

1451 Manuscript submitted to ACM

1452

- [11] Sarah Cohen, Werner Nutt, and Yehoshua Sagiv. 2007. Deciding equivalences among conjunctive aggregate queries. *J. ACM* 54, 2, Article 5 (April 2007), 50 pages. [doi:10.1145/1219092.1219093](https://doi.org/10.1145/1219092.1219093)
- [12] Bruce P. Douglass, David Harel, and Mark B. Trakhtenbrot. 1998. Statecharts in use: structured analysis and object-orientation. In *Lectures on Embedded Systems*, Grzegorz Rozenberg and Frits W. Vaandrager (Eds.). Lecture Notes in Computer Science, Vol. 1494. Springer-Verlag, London, 368–394. [doi:10.1007/3-540-65193-4_29](https://doi.org/10.1007/3-540-65193-4_29)
- [13] Ian Editor (Ed.). 2007. *The title of book one* (1st. ed.). The name of the series one, Vol. 9. University of Chicago Press, Chicago, Chapter The title of the chapter, 127–238. [doi:10.1007/3-540-09237-4](https://doi.org/10.1007/3-540-09237-4)
- [14] Ian Editor (Ed.). 2008. *The title of book two* (2nd. ed.). University of Chicago Press, Chicago, Chapter 100, 25–137. [doi:10.1007/3-540-09237-4](https://doi.org/10.1007/3-540-09237-4)
- [15] Matthew Van Gundy, Davide Balzarotti, and Giovanni Vigna. 2007. Catch me, if you can: Evading network signatures with web-based polymorphic worms. In *Proceedings of the first USENIX workshop on Offensive Technologies (WOOT '07)*. USENIX Association, Berkley, CA, Article 7, 9 pages.
- [16] Torben Hagerup, Kurt Mehlhorn, and J. Ian Munro. 1993. Maintaining Discrete Probability Distributions Optimally. In *Proceedings of the 20th International Colloquium on Automata, Languages and Programming (Lecture Notes in Computer Science, Vol. 700)*. Springer-Verlag, Berlin, 253–264.
- [17] David Harel. 1978. *LOGICS of Programs: AXIOMATICS and DESCRIPTIVE POWER*. MIT Research Lab Technical Report TR-200. Massachusetts Institute of Technology, Cambridge, MA.
- [18] David Harel. 1979. *First-Order Dynamic Logic*. Lecture Notes in Computer Science, Vol. 68. Springer-Verlag, New York, NY. [doi:10.1007/3-540-09237-4](https://doi.org/10.1007/3-540-09237-4)
- [19] Lars Hörmander. 1985. *The analysis of linear partial differential operators. III*. Grundlehren der Mathematischen Wissenschaften [Fundamental Principles of Mathematical Sciences], Vol. 275. Springer-Verlag, Berlin, Germany. viii+525 pages. Pseudodifferential operators.
- [20] Lars Hörmander. 1985. *The analysis of linear partial differential operators. IV*. Grundlehren der Mathematischen Wissenschaften [Fundamental Principles of Mathematical Sciences], Vol. 275. Springer-Verlag, Berlin, Germany. vii+352 pages. Fourier integral operators.
- [21] IEEE 2004. IEEE TCSC Executive Committee. In *Proceedings of the IEEE International Conference on Web Services (ICWS '04)*. IEEE Computer Society, Washington, DC, USA, 21–22. [doi:10.1109/ICWS.2004.64](https://doi.org/10.1109/ICWS.2004.64)
- [22] Markus Kirschmer and John Voight. 2010. Algorithmic Enumeration of Ideal Classes for Quaternion Orders. *SIAM J. Comput.* 39, 5 (Jan. 2010), 1714–1747. [doi:10.1137/080734467](https://doi.org/10.1137/080734467)
- [23] Donald E. Knuth. 1997. *The Art of Computer Programming, Vol. 1: Fundamental Algorithms* (3rd. ed.). Addison Wesley Longman Publishing Co., Inc., Boston.
- [24] David Kosiur. 2001. *Understanding Policy-Based Networking* (2nd. ed.). Wiley, New York, NY.
- [25] Newton Lee. 2005. Interview with Bill Kinder: January 13, 2005. Video. *Comput. Entertain.* 3, 1, Article 4 (Jan.-March 2005). [doi:10.1145/1057270.1057278](https://doi.org/10.1145/1057270.1057278)
- [26] Dave Novak. 2003. Solder man. Video. In *ACM SIGGRAPH 2003 Video Review on Animation theater Program: Part I - Vol. 145 (July 27–27, 2003)*. ACM Press, New York, NY, 4. [doi:10.945/woot07-S422](https://doi.org/10.945/woot07-S422) <http://video.google.com/videoplay?docid=6528042696351994555>
- [27] Barack Obama. 2008. A more perfect union. Video. Retrieved March 21, 2008 from <http://video.google.com/videoplay?docid=6528042696351994555>
- [28] Poker-Edge.Com. 2006. Stats and Analysis. Retrieved June 7, 2006 from <http://www.poker-edge.com/stats.php>
- [29] R Core Team. 2019. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria. <https://www.R-project.org/>
- [30] Brian J. Reiser. 2014. Designing coherent storylines aligned with NGSS for the K-12 classroom. Presentation at National Science Education Leadership Association Meeting, Boston, MA, USA. <https://www.academia.edu/6884962/>
- [31] Bernard Rous. 2008. The Enabling of Digital Libraries. *Digital Libraries* 12, 3, Article 5 (July 2008). To appear.
- [32] Mehdi Saeedi, Morteza Saheb Zamani, and Mehdi Sedighi. 2010. A library-based synthesis methodology for reversible logic. *Microelectron. J.* 41, 4 (April 2010), 185–194.
- [33] Mehdi Saeedi, Morteza Saheb Zamani, Mehdi Sedighi, and Zahra Sasanian. 2010. Synthesis of Reversible Circuit Using Cycle-Based Approach. *J. Emerg. Technol. Comput. Syst.* 6, 4 (Dec. 2010), 12 pages.
- [34] Joseph Scientist. 2009. The fountain of youth. Patent No. 12345, Filed July 1st., 2008, Issued Aug. 9th., 2009.
- [35] Stan W. Smith. 2010. An experiment in bibliographic mark-up: Parsing metadata for XML export. In *Proceedings of the 3rd. annual workshop on Librarians and Computers (LAC '10, Vol. 3)*, Reginald N. Smythe and Alexander Noble (Eds.). Paparazzi Press, Milan Italy, 422–431.
- [36] Asad Z. Spector. 1990. Achieving application requirements. In *Distributed Systems* (2nd. ed.), Sape Mullender (Ed.). ACM Press, New York, NY, 19–33. [doi:10.1145/90417.90738](https://doi.org/10.1145/90417.90738)
- [37] Harry Thornburg. 2001. *Introduction to Bayesian Statistics*. Retrieved March 2, 2005 from <http://ccrma.stanford.edu/~jos/bayes/bayes.html>, archived at <https://web.archive.org/web/20240505055615/https://ccrma.stanford.edu/~jos/bayes/bayes.html>
- [38] TUG 2017. *Institutional members of the TeX Users Group*. Retrieved May 27, 2017 from <http://www.tug.org/instmem.html>
- [39] Boris Veysman. 2017. *acmart—Class for typesetting publications of ACM*. Retrieved May 27, 2017 from <http://www.ctan.org/pkg/acmart>

Received 20 February 2007; revised 12 March 2009; accepted 5 June 2009